

Transport layer. TCP, UDP.

Ing. Yelena Trofimova

Department of Computer Systems
Faculty of Information Technology
Czech Technical University in Prague
©Yelena Trofimova, 2021

Computer networks, BIE-PSI
SS 2020/21, Lecture 6

<https://courses.fit.cvut.cz/BIE-PSI/>



Content

- transport layer
 - ▶ layer services
 - ▶ protocols
 - ▶ establishment and termination of connection
 - ▶ flow control, congestion control
- TCP
- UDP
- RTP, RTCP

Transport layer services

- transport layer is a boundary between
 - ▶ applications (software)
and
 - ▶ networks (technology)
- ensures transparent data transfer between end-users
 - ▶ reliability
 - ▶ flow control
 - ▶ data segmentation
 - ▶ correction of errors

Addressing

- service must have an address
 - ▶ TSAP – Transport Service Access Point
- predefined TSAPs ("well-known")
 - ▶ eg. port 80 for http
- dynamically allocated – needs special software
 - ▶ portmapper
 - registers TSAP for services
 - eg. BitTorrent

Connection establishment I

- trivial solution:

CONNECTION REQUEST — — — >
< — — — CONNECTION ACCEPTED

problem: duplicated CR packet

- improvement:
 - ▶ unique identifier for each session
 - ▶ limited lifetime of packet

Connection establishment II

- three-way handshake:

CONNECTION REQUEST — — — >
(seq=x)

< — — — ACKNOWLEDGE
(seq=y, ack=x)


DATA — — — >
(seq=x+1, ack=y)

Connection termination

- asymmetric termination
 - ▶ one user terminates the connection (example: landline phone call)
 - ▶ possible data loss when the other user sent data before receiving termination message
- symmetric termination
 - ▶ both sides must agree
 - ▶ there is no fully reliable solution if communication is not secured (Two Armies Problem / Coordinated Attack Problem / Two Generals' Problem)

Two Armies Problem

A will be winner if A1 and A2 start attack at one moment.



- A1 sends the message, but it is not clear if the message passes the enemy area
- A2 sends the acknowledgment, but it can be also lost
- A1 should send the acknowledgment of received acknowledgment
- It is still not clear if they agreed: another acknowledgment?
- Prove of non-existence of solution is by contradiction: let n be the number of messages enough for the protocol to work correctly. If the last message is lost: sender of the last message will attack and the other side will not – contradiction.

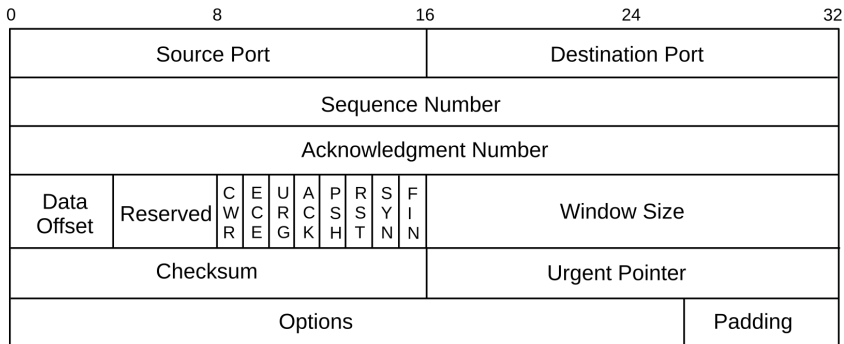
Flow control and congestion control

- similar principle as at the link layer
 - ▶ detection of errors
 - ▶ elimination of duplicated packets (sequence numbers)
 - ▶ limited number of unacknowledged packets
 - ▶ sliding window
- congestion control
 - ▶ prevents from overwhelming the network

TCP


- Transmission Control Protocol
- service on layer 4
 - ▶ connection-oriented
 - ▶ guarantees error free data transmission
 - ▶ duplex
- many implementations – various improvements
 - ▶ Reno
 - ▶ Tahoe
 - ▶ Vegas
 - ▶ New Reno, CUBIC and many others

TCP header



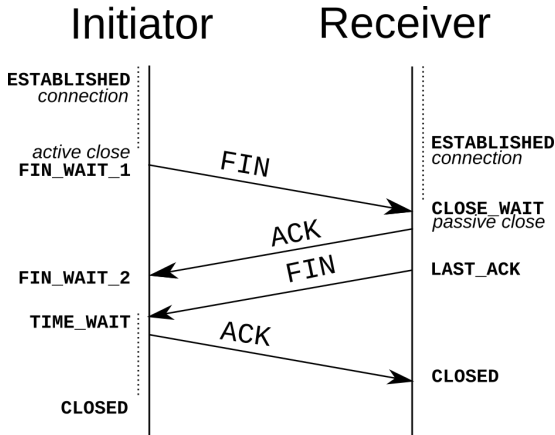
Source: <https://intronetworks.cs.luc.edu/current/html/tcp.html>

TCP connection establishment



Source: <https://lwn.net/Articles/508865/>

TCP connection termination



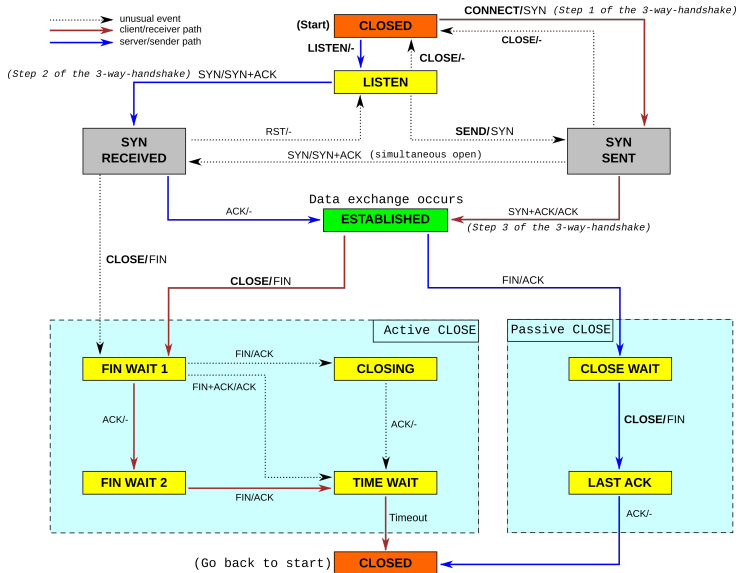
Source: Wikipedia

Note: after second message (ACK from receiver) connection is half-closed and receiver can still send data to initiator.

TCP properties

- data delivery guarantee:
 - ▶ checksum
 - ▶ detection of duplicate packets
 - ▶ retransmission
 - ▶ correct ordering
 - ▶ timeout
- sliding window
- congestion detection

TCP state diagram (source: Wikipedia)



Sliding window

- receiver reserves buffer – window for receiving data
- receiver specifies the current window size in the reply
 - ▶ sender must not send more data than the size of the window, then waits for acknowledgment
 - ▶ window size is decreased if the receiver does not manage to process the data
 - ▶ the window and its shrinking are used for flow control

Acknowledgment I

- receiver sends ACK, where is the sequence number of the byte, which is expected
- this confirms that all the previous bytes has been received
- receiver does not need to send acknowledgment individually for all packets

Timeout (on the sender side):

- sender has no ACK up to expected time
- unacknowledged data will be retransmitted

Acknowledgment II

Duplicate ACK

- if a packet does not arrive, but the following does, receiver repeats the last ACK
- it will be used as an indicator that the packet might have been lost

Timeout (on receiver side)

- no packet arrived up to expected time
- send again the last ACK
- max. 3 retransmissions of ACK

Window size

- 16 bits \Rightarrow max 64kB
- problem for fast links or great delay
 - ▶ bandwidth * delay
 - ▶ example: 1Gbps link:
 - sending of 64kB takes 0.5 ms
 - for delay of 50 ms link is used only for 1% of its capacity
- solution: window scale option
 - ▶ a maximum value of 14 may be used for the scale factor \Rightarrow 30 bits for window size

Congestion control variables

- MSS (Maximum Segment Size) – maximum size of the packet to be sent, defined by the receiver
- SSTHRESH (Slow-start Threshold) – limit for likely congestion
 - ▶ estimate how many unconfirmed data can be sent
 - ▶ is multiples of MSS
- CWND (Congestion Window) – sender window
 - ▶ how many unconfirmed data the sender sends

Congestion Control I


- limit on size of unacknowledged bytes
- slow start: initial value of CWND is 1
- congestion window size is increased 2x: grows exponentially
- after reaching SSTHRESH size is grown linearly

Congestion Control II

Congestion:

- packet or its acknowledgment is lost
- if ACK did not come at all (timeout):
 - ▶ transmission is repeated from the last acknowledged packet
 - ▶ slow start is repeated
- is duplicate ACK comes 3x:
 - ▶ "fast retransmit" and "fast recovery" is activated
 - ▶ lost packet is sent again, but not the following ones ("selective retransmit")
 - ▶ CWND is decreased by half

Evolution of TCP's congestion window (Tahoe and Reno)



Source: Kurose & Ross

TCP usage

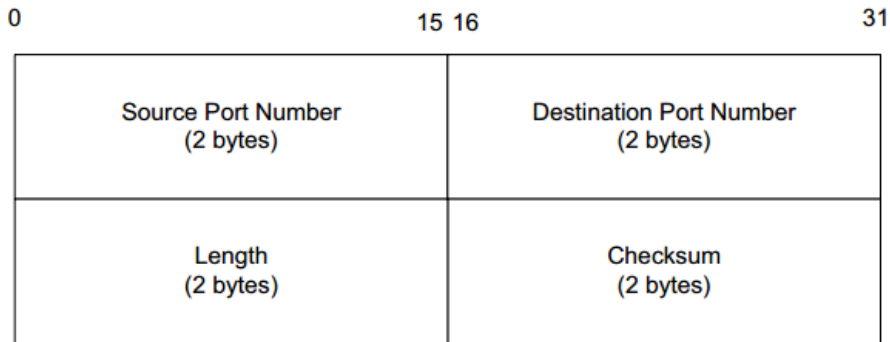
- TCP is good for applications which need reliable data channel
- TCP is not good for small blocks of data – high overhead (time and data)
- TCP is not suitable for real-time applications (VOIP, streaming)
 - ▶ the packet is delivered at all costs – unacceptable delays
- TCP is not suitable for embedded-systems (too complex)

UDP

User Datagram Protocol

- service on layer 4
- connection-less
- UDP datagram may be lost
- uses ports for multiplexing as TCP
- maximum 64 kB of data (avoid fragmentation in IP)
- example: DNS

UDP header



UDP usage

- small blocks of data, when:
 - ▶ packet losses are not be critical (eg. DNS)
 - ▶ TCP overhead if not acceptable (eg. NTP)
- real-time applications

RTP

Real-time Transport Protocol

- streaming data between endpoints in real time
- header contains "timestamp"
 - ▶ time from the beginning of the stream
 - ▶ the unit depends on the application
 - ▶ allows to interpret the received block
- implemented mostly over UDP
- RTP usage:
 - ▶ any multimedia formats (H.264, MPEG-4, MJPEG, MPEG, ...)
 - ▶ videoconferencing
 - ▶ data streaming
 - ▶ IP Telephony

RTCP

RTP Control Protocol

- provides statistics and control information for an RTP session
- does not transport any media data itself
- used to control quality of service