

Web Data Mining

Lecture 2: Data Access and Acquisition Methods

Milan Dojčinovski
milan.dojchinovski@fit.cvut.cz



Czech Technical University in Prague - Faculty of Information Technologies - Software and Web Engineering



Summer semester 2019/2020
Humla v0.3

Overview

- **Introduction**
- Web Crawler Architecture
- Crawl Control Technologies

Crawling in a Nutshell

- Automatic harvesting of web content
 - *texts, images, videos, tweets, pdfs, spreadsheets, etc.*
- Specialized intelligent programs - Web Crawlers
 - *also known as robots, bots or spiders*
- Recursive visit of web pages
 - *visit a seed web page URL, download content, extract links, add new links to the queue*
- Application of set of policies (rules)
 - *do not visit already visited pages*
 - *ignore links to images, videos or other links not pointing to a web page*

Challenges


- The Web is big
 - *insufficient resources to crawl entire web*
 - *only valuable should be visited, downloaded, indexed*
- Ethical issues
 - *do not overload web servers with requests*
 - *the crawlers should identify themselves*
 - *compliance with the Robot Exclusion Protocol*
- Keep crawled content fresh
 - *check content update and perform incremental updates*
- Detect "search engine spamming"
 - *tricking the ranking algorithms to achieve higher rank in search results*

Applications

- Web Search Engines
 - Google, Bing, ...
- Web Archiving
 - Digital preservation, archives, ...
- Vertical Search Engines
 - Car rentals, apartments, ...
- Web Data Mining
 - Focused crawlers, ...
- Web monitoring
- Malicious web sites detections, searching for illegal content, ...
- Web site/applications testing
- Mirroring

Bot Traffic Report

- Good: 12.2% Feed fetchers, 6.6% Search engine bots, 2.9% Commercial crawlers, 1,2% Monitoring bots
- Bad: 24.3% Impersonators, 2.6% Hacker tools, 1.7% Scrapers, 0.3% Spammers
- <https://www.incapsula.com/blog/bot-traffic-report-2016.html>

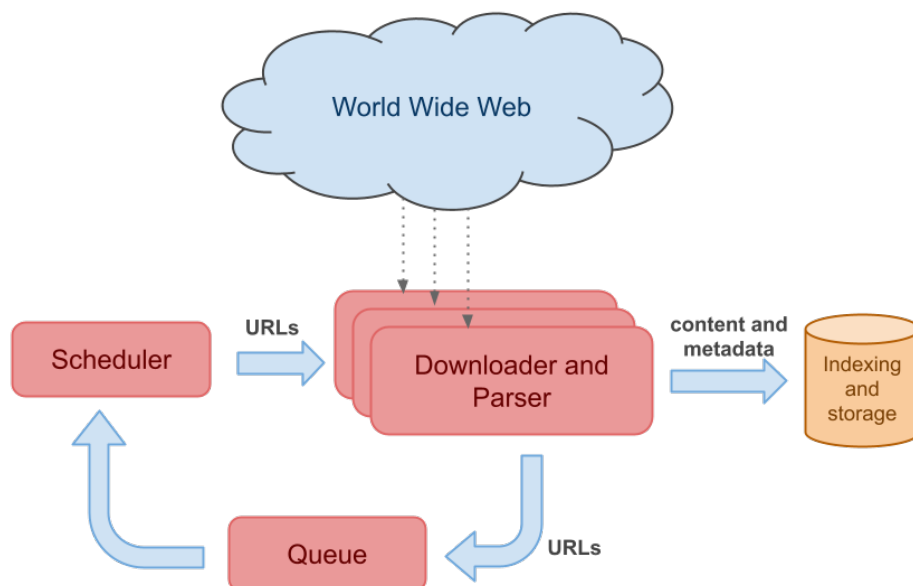


Overview

- Introduction
- **Web Crawler Architecture**
- Crawl Control Technologies

High-level crawler architecture

- A standard web crawler



Crawler Types

- Batch Crawlers
 - *create a snapshot of a web space*
- Incremental Crawlers
 - *continuously crawl a web space*
 - *revisit URLs to ensure freshness*
- Focused Crawlers
 - *crawl only pages in certain categories*
 - *e.g. pages from only specific domain, with high popularity, specific content type, ...*
- Topical Crawlers
 - *crawl only pages in certain topics*
 - *e.g. sports, real estate, ...*

Terminology

- Seed pages
 - *a list of URLs to visit first*
- Frontier
 - *holds a list of unvisited URLs*
- Fetcher
 - *fetches the contents of a web page*
 - *usually for the data transfer is used the HTTP protocol*
- Link Extractor
 - *parses the HTML contents and extracts URLs to other pages*
- Other Components
 - *URL filter*
 - *checks whether URL is on the black-list*
 - *Duplicate eliminator*
 - *eliminates already visited URLs*
 - *URL prioritizer*
 - *assigns priority to a new URL*
 - *only internal(external) url internal first specific language*

A Basic Crawling Algorithm

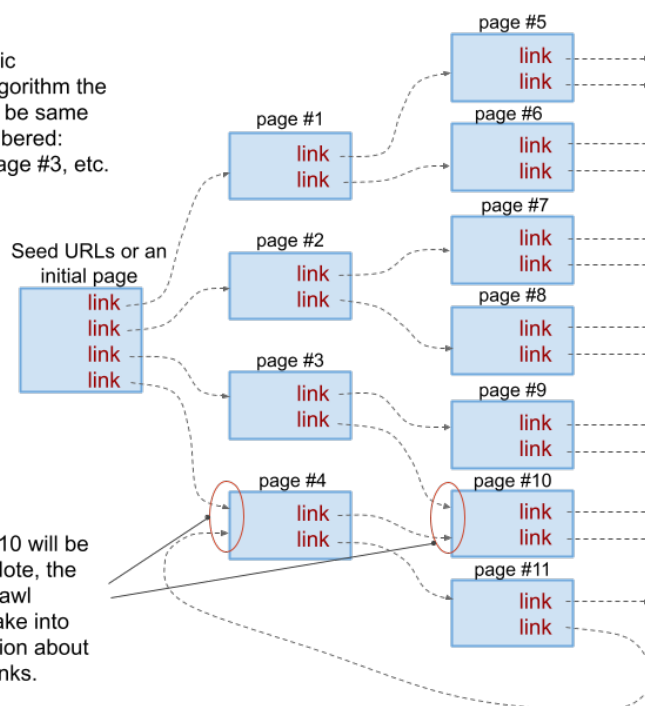
- A breadth-first crawl algorithm:
 - *Input*
 - `url_queue` - seed pages
 - *Data structures used:*
 - `crawled_pages` - list of already visited pages
- Steps:

```
1 while url_queue is not empty:
2     url = get the head (first) element from the url_queue
3     # The Frontier maintains the url_queue
4     page = fetch the page from url
5     new_links = extract links from the page's contents
6     for each link in new_links:
7         if link is does not exist in url_queue AND crawled_pages
8             add it at the end of the url_queue list
```

- Order in which pages are visited
 - *is highly correlated with their popularity*
 - *bias of search engines to index well connected pages*

A Breadth-First Crawling

- According to the basic breadth-first crawl algorithm the order of crawling will be same as the page are numbered: page #1, page #2, page #3, etc.



- Page #4 and page #10 will be crawled only once. Note, the basic breadth-first crawl algorithm does not take into account the information about the number of backlinks.

Backlink Based Crawling Strategy

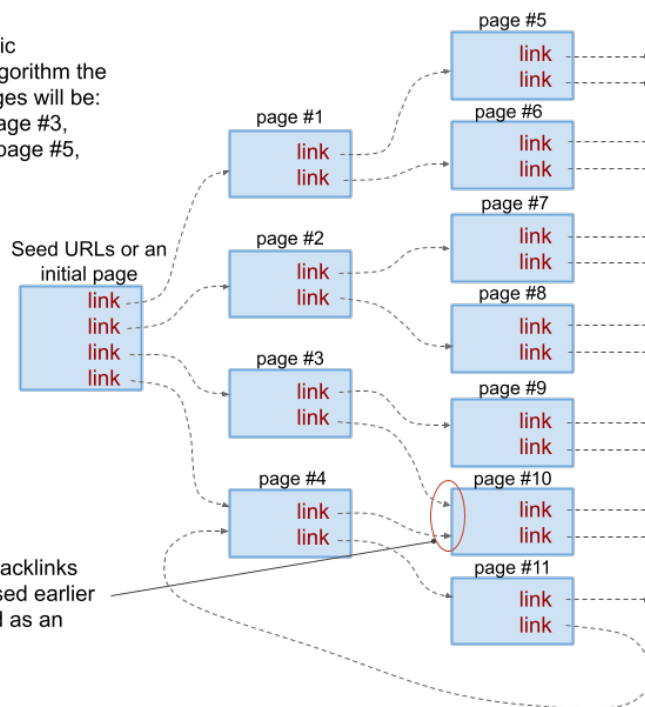
- How to select the next URL for crawling?
- Sort the URLs in the **url_queue** according to their relevance
 - take into account the backlinks
 - backlink
 - an incoming link to a web page
- Function for sorting of URLs in the **url_queue**. Steps:

```
1 | foreach url in url_queue:  
2 |     # count the number of links pointing to url  
3 |     backlink_count[url]  
4 |     sort url_queue by the backlink_count[url] information
```

- URLs with more incoming links will be processed first!
- Sorting of the **url_queue** is done after adding new links from a page!

Backlink Based Crawling

- According to the basic breadth-first crawl algorithm the order of crawling pages will be: page #1, page #2, page #3, page #4, **page #10**, page #5, page #6, ...



- Page #10 has two backlinks and it will be processed earlier since it is considered as an important page.

Fetching

- The Crawler acts as a Web client
- Sends an HTTP request to the server hosting the page and receives the response
- The crawler sends HTTP request to the server
 - Resolving the host name in the URL to an IP address using DNS
 - Connecting a socket to the server and sending the request

```
1 GET /wiki/Jan_Palach HTTP/1.1
2 Host: en.wikipedia.org
3 Accept: text/html
```

- The server sends response with the content of the resource
http://en.wikipedia.org/wiki/Jan_Palach

```
1 HTTP/1.1 200 OK
2 Content-Type: text/html; charset=UTF-8
3 [the other headers are not shown due to brevity]
4
5 <!DOCTYPE HTML>
6 <head>
7   <meta charset="UTF-8" />
8   <title>Jan Palach - Wikipedia, the free encyclopedia
9 </head>
10 <body>
11 ...
12 </body>
```

Parsing

- Parsing the content of the HTTP payload
 - *extracting content for indexing*
 - *extracting links to be added to the frontier*
 - *extracting additional crawling and indexing directives*
 - *headers Cache-Control, Content-Type, X-Robots-Tag, ...*
- HTML code very often contains invalid markup
 - *unclosed elements, unencoded special characters, missing required attributes, improperly nested tags, missing quotes, ...*
- Bad HTML markup should be fixed
 - *a preprocessing step is required to clean up the HTML*
 - *many tools available, tidy - a tool provided by W3C*

URL Normalization

- Transform a URL into its normalized or canonical form
- URL Normalization is part of the
 - *W3C RFC 3986 specification*
- Relative URLs that are in a parsed document need normalization
 - e.g., the website http://en.wikipedia.org/wiki/Jan_Palach contains a relative URL: /wiki/Prague_Spring
 - *fully qualified domain name and path are not specified*
- Normalization of relative URLs
 - *relative URLs should be expanded with the fully qualified domain*
 - /wiki/Prague_Spring becomes http://en.wikipedia.org/wiki/Prague_Spring
 - *they do not contain the fully qualified domain name and path*

URL Normalization (cont.)

- Case Normalization
 - *the URI scheme and the protocol should be normalized to lowercase*
 - <hTtp://www.EXAMPLE.com/> is equivalent to <http://www.example.com/>
 - *hexadecimal digits within a percent-encoding triplet should be normalized to uppercase (for the digits A-F)*
 - e.g., ["%3a"](#) to ["%3A"](#)
- Percent-Encoding Normalization
 - *decoding any percent-encoded octet*
 - <http://www.example.com/%7Eusername/> becomes <http://www.example.com/~username/>
 - *digits, alphabet letters, hyphens (-), period (.), underscores (_), tildes (~) should not be encoded in URIs*
- Path Segment Normalization
 - *removing dot-segments "." and ".." according to the algorithm described in RFC 3986 specs*
 - <http://example.com/a/b/./c/> becomes <http://example.com/a/c/>

URL Normalization (cont.)

- Scheme-Based Normalization
 - *may vary from scheme to scheme*
 - *for the "http" scheme, the default port "80" can be removed*
 - *http://example.com:80/ becomes http://example.com/*
- Others:
 - *Replacing IP with domain name*
 - *Removing the fragment*
 - *Sorting the query parameters*
 - *Removing unused query variables*
 - ...
- Upon performed normalization, we can use the absolute URL

Crawler Identification

- Keep track of the crawling
- Using User-Agent HTTP header
- Example:
 - **User-Agent: Googlebot**
 - *an user-agent string which identifies the Google's crawler*
 - *list of user-agent strings used by Google's crawlers*
- Crawler developers might consider include also additional info about the crawler
 - *e.g., URL of crawler website*

Revisit Strategies

- Pages are being added, updated or deleted continuously.
 - *Need to revisit pages to avoid outdated data.*
 - *Older a page gets, the more it costs not to crawl it.*
 - *Not possible to regularly check all pages!*
- Strategies
 - *Uniform revisiting*
 - *Revisit all regularly*
 - *Proportional revisiting*
 - *Revisit more frequently more frequently changing pages*
 - *Hybrid solution*
 - *Balance resources and frequency of changes*

Revisit Strategies (cont.)

- Crawlers can use HTTP HEAD to check metadata (e.g. Last-Modified, Expires headers)
 - *Freshness*
 - *The proportion of pages that are fresh*
 - *Can lead to bad decisions, such as not crawling popular sites*
 - *Age*
 - *Is the # of "days" that an average page is out-of-date.*

$$F_p(t) = \begin{cases} 1 & \text{if } p \text{ is equal to the local copy at time } t \\ 0 & \text{otherwise} \end{cases}$$

$$A_p(t) = \begin{cases} 0 & \text{if } p \text{ is not modified at time } t \\ t - \text{mod. time of } p & \text{otherwise} \end{cases}$$

Scrapy Example

```
1 import scrapy
2 class MySpider(scrapy.Spider):
3     name = 'novinky.cz'
4     start_urls = ['https://www.novinky.cz']
5
6     def parse(self, response):
7         for title in response.css('div.item > h3'):
8             yield {'title': title.css('a ::text').extract_first()}
9
10        for next_page in response.css('div.menu > ul > li'):
11            yield scrapy.Request(
12                response.urljoin(next_page.css('::attr(href)').extract_firs
13                callback=self.parse)
```

```
1 | scrapy runspider myspider.py
```

```
1 { 'downloader/request_bytes': 7303,
2   'downloader/request_count': 29,
3   'downloader/request_method_count/GET': 29,
4   'downloader/response_bytes': 1045117,
5   'downloader/response_count': 29,
6   'downloader/response_status_count/200': 28,
7   'downloader/response_status_count/301': 1,
8   'dupefilter/filtered': 1188,
9   'finish_reason': 'finished',
10  'finish_time': datetime.datetime(2017, 2, 9, 14, 6, 13, 928444),
11  'item_scraped_count': 196,
12  'log_count/DEBUG': 227,
```

Lecture 2: Data Access and Acquisition Methods - Milan Dojčinovski

- 23 -

Overview

- Introduction
- Web Crawler Architecture
- Crawl Control Technologies

Lecture 2: Data Access and Acquisition Methods - Milan Dojčinovski

- 24 -

Robot Exclusion Protocol

- A de-facto standard defining etiquette policies
 - <http://www.robotstxt.org/orig.html>
 - Consensus from 1994 between the majority of robot authors and other people with an interest in robots.
 - Not owned by any standardization body like W3C
- Specifies access policies/instructions about their sites to the web robots
 - which sites not to crawl, which bots
- Its a simple text file: **robots.txt**
 - Accessible via HTTP on the local URL **"/robots.txt"**

Robots.txt Fields

- Format
 - defined as text UTF-8 encoded file named **robots.txt**
 - must be accessible via HTTP and present in the top-level directory
 - one or more records separated by one or more blank lines
 - each record contains lines of the form:
- ```
1 | <field>:<optionalspace><value><optionalspace>
```
- A record starts with one or more **User-agent** lines, followed by one or more **Disallow** lines
  - **User-agent** field
    - the name of the robot the record is describing access policy for
    - if more User-agent fields are present, the same policies apply to all robots
    - at least one User-agent field should be present
  - **Disallow** field
    - a partial URL that is not to be visited
    - any URL that starts with this value will not be retrieved
    - empty value indicates that all URLs can be retrieved

## Robots.txt Examples

- Allow complete access

```
1 | User-agent: *
2 | Disallow:
```

- Go away!

```
1 | User-agent: *
2 | Disallow: /
```

- Specific robot

```
1 | User-agent: GoogleBot
2 | Disallow: /misc/ # Here are scripts. Do not index scripts (JavaScript).
3 | # This matches /misc/script.js, but also /misc/a/script.js
4 |
5 | # Files
6 | Disallow: /CHANGELOG.txt # Do not index the changelog file.
7 |
8 | # Paths (clean URLs)
9 | Disallow: /tmp/ # Files in /tmp/ will soon disappear.
10 |
11 | # Paths (no clean URLs)
12 | Disallow: /?q=logout/ # Do we really need to index the logout page?
```

- Add every URL or set of URLs that you want to exclude indexing

## Robot Exclusion Protocol

- Robots.txt File Example  
– excerpt from the robots.txt at <http://fit.cvut.cz/robots.txt>

```
1 | User-agent: *
2 | Crawl-delay: 10
3 | # Directories
4 | Disallow: /includes/
5 | Disallow: /misc/
6 | Disallow: /modules/
7 | # Files
8 | Disallow: /CHANGELOG.txt
9 | Disallow: /cron.php
10 | Disallow: /INSTALL.mysql.txt
11 | Disallow: /INSTALL.pgsql.txt
12 | Disallow: /install.php
13 | Disallow: /INSTALL.txt
14 | # Paths (clean URLs)
15 | Disallow: /admin/
16 | Disallow: /comment/reply/
17 | Disallow: /filter/tips/
18 | Disallow: /logout/
19 | Disallow: /node/add/
20 | # Paths (no clean URLs)
21 | Disallow: /?q=admin/
22 | Disallow: /?q=comment/reply/
23 | Disallow: /?q=filter/tips/
24 | Disallow: /?q=logout/
```

## Robots.txt Extensions

- **Crawl-delay** field
  - e.g. number of seconds to wait between subsequent visits
- **Allow** field
  - specifies paths that may be accessed by the designated crawlers
  - useful when one tells robots to avoid an entire directory but still wants some HTML documents in that directory crawled
  - supported by Google, Microsoft (Bing), Yahoo
- **Sitemap** field
  - URL pointer to the Sitemap index file
  - the URL should be absolute!
  - empty value indicates that all URLs can be retrieved
- **Host** field
  - Specify preferred domain for websites with multiple mirrors

```
1 User-agent: *
2 Allow: /
3 Host: www.example.com
4 Sitemap: http://www.example.com/sitemap.xml
```

## Page-level Indexing Configuration

- Control how robots make content available through search results
  - including a meta tag element in the HTML page, or
  - an HTTP response header
- Robots meta tag

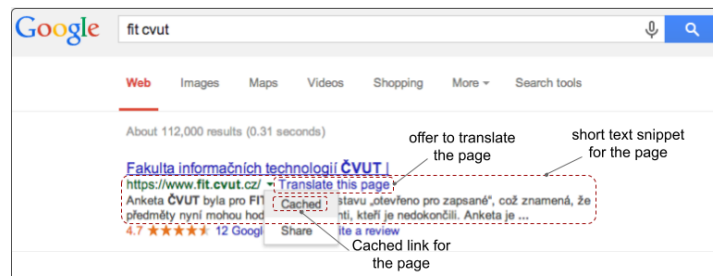
```
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <meta name="robots" content="noindex" />
5 (...)
6 </head>
7 <body> (...) </body>
8 </html>
```

- Instructs all search engines not to show the page in search

```
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <meta name="googlebot" content="noindex" />
5 (...)
6 </head>
7 <body> (...) </body>
8 </html>
```

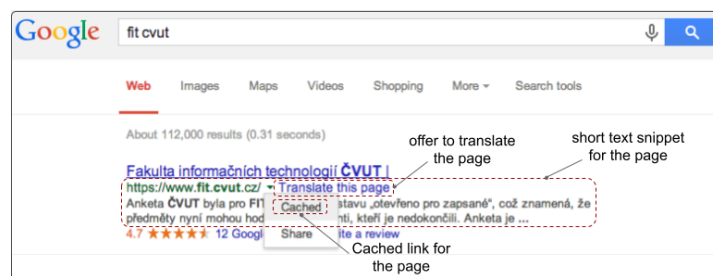
- Instructs Google not to show this page in its search results

# Indexing and Serving Directives



- **noindex**: do not show this page in search results, do not show cached link
- **noarchive**: do not show a "Cached" link in search results
- **none**: same as noindex, nofollow
- **nosnippet**: do not show a snippet in the search results for this page
- **notranslate**: do not offer translation of this page in search results

# Indexing and Serving Directives (cont.)



- **all**: no restrictions for indexing or serving
- **nofollow**: do not follow the links on this page
- **noimageindex**: do not index images on this page
- **unavailable\_after**: [RFC-850 date/time]: do not show this page in search results after the specified date/time
- **noodp**: don't use metadata (titles or snippets) from the Open Directory project (closed on March 17, 2017)



## X-Robots-Tag HTTP Header

- HTTP header, part of the HTTP response
- Any directive from the robots meta tag can be specified
- Usage example
  - *HTTP request-response*
  - **X-Robots-Tag: noindex** - *the server instructs the crawler not to index the page*

```
1 > GET /index.html HTTP/1.1
2 > Host: example.com
3 > Accept: text/html
4 > Date: Thu, 27 Feb 2014 10:12:77 GMT
5
6 < HTTP/1.1 200 OK
7 < Content-Type: text/html
8 < X-Robots-Tag: noindex
9 <
10 < ...resource representation in html...
```

- Bot specification
  - **X-Robots-Tag: googlebot: noindex**

## X-Robots-Tag HTTP Header

- Multiple **X-Robots-Tag** headers can be specified in the HTTP response
- Multiple headers example. Specified directives:
  - *no indexing of images in the requested page*
  - *the page will not be available after April 22nd 2014 23:59:59*

```
1 > GET /index.html HTTP/1.1
2 > Host: example.com
3 > Accept: text/html
4 > Date: Thu, 27 Feb 2014 10:12:77 GMT
5
6 < HTTP/1.1 200 OK
7 < Content-Type: text/html
8 < X-Robots-Tag: noimageindex
9 < X-Robots-Tag: unavailable_after: 22 Apr 2014 23:59:59 PST
10 <
11 < ...resource representation in html...
```

- multiple directives can be also specified in one single header - comma-separated

```
1 < X-Robots-Tag: noimageindex, unavailable_after: 22 Apr 2014 23:59:59 PST
```

## Web Servers Configuration for X-Robots-Tag

- Apache **X-Robots-Tag** header configuration
  - *Configuring Apache with **noindex** and **nofollow** directives*
  - *Do not index JavaScript files across the entire web site*

```
1 <Files ~ "\.js$">
2 Header set X-Robots-Tag "noindex, nofollow"
3 </Files>
```

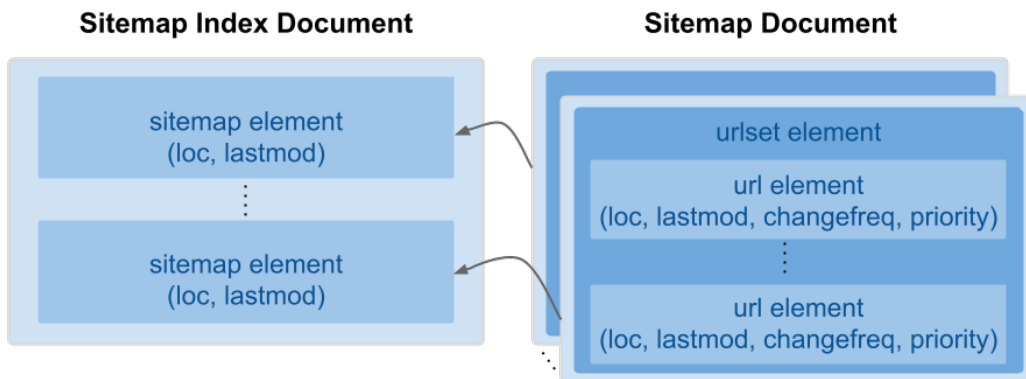
- Node.js **X-Robots-Tag** header configuration
  - *Configuring Node.js with **noindex** directive*

```
1 var http = require('http');
2
3 http.createServer(function (req, res) {
4 // assuming variable 'file' holds the contents of the file
5 // setting the X-Robots-Tag response header
6 res.setHeader('X-Robots-Tag', 'noindex');
7 res.end(file);
8 }).listen(8080);
```

## Sitemaps

- Need for standard web site structure format
  - *Machine-processable Web site structure description*
  - *Complementary work to **robots.txt***
    - **robots.txt** spec does not say how to instruct crawlers what to crawl, how often, etc.
- Standard protocol describing web site structure
  - <https://www.sitemaps.org>
  - *URLs available for crawling*
- Introduced by Google in June 2005, sitemaps main site
- Supported by Google, Yahoo! and Microsoft
- An XML file with list of URLs and URLs metadata
- Sitemaps is a robot inclusion standard

## Sitemaps Format



- Two types of sitemap documents
  - *Sitemap Index Document*
    - provides aggregation of multiple sitemap files, together with their location and their last modification timestamp
  - *Sitemap Document*
    - represent a single sitemap document containing a list of URLs offered for crawling

## Sitemap Format

- Sitemap Index Document Example
  - excerpt from the **sitemap.xml** at <http://fit.cvut.cz/sitemap.xml>
  - groups two sitemap files

```
1 <sitemapindex xmlns="http://www.sitemaps.org/schemas/sitemap/0.9"
2 xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
3 xsi:schemaLocation="http://www.sitemaps.org/schemas/sitemap/0.9 http://w
4 <sitemap>
5 <loc>http://fit.cvut.cz/sitemap0.xml</loc>
6 <lastmod>2014-02-27T15:26:35+00:00</lastmod>
7 </sitemap>
8 <sitemap>
9 <loc>http://fit.cvut.cz/sitemap1.xml</loc>
10 <lastmod>2010-02-15T15:00:56+00:00</lastmod>
11 </sitemap>
12 </sitemapindex>
```

## Sitemap Example

- Sitemap Format Example
  - excerpt from the `sitemap0.xml` at <http://fit.cvut.cz/sitemap0.xml>
  - an sitemap file
  - groups multiple URLs

```
1 <urlset xmlns="http://www.sitemaps.org/schemas/sitemap/0.9"
2 xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
3 xsi:schemaLocation="http://www.sitemaps.org/schemas/sitemap/0.9 http://w
4 <url>
5 <loc>http://fit.cvut.cz/</loc>
6 <lastmod>2014-02-27T15:26:35+00:00</lastmod>
7 <changefreq>hourly</changefreq>
8 <priority>1.0</priority>
9 </url>
10 <url>
11 <loc>http://fit.cvut.cz/node/1</loc>
12 <lastmod>2010-07-31T08:23:17+00:00</lastmod>
13 <changefreq>always</changefreq>
14 <priority>0.8</priority>
15 </url>
16 ...
17 </urlset>
```

## Other sources of metadata

- Well-Known URIs
  - site-wide metadata
  - location `/.well-known/`
  - Examples
    - `/.well-known/host-meta`
      - host metadata including author, copyright, etc.
    - `/.well-known/security.txt`
      - contacts for reporting security vulnerabilities, ...