

Web Data Mining

Lecture 10: Recommender Systems

Jaroslav Kuchař & Milan Dojčinovski

jaroslav.kuchar@fit.cvut.cz, milan.dojchinovski@fit.cvut.cz



Czech Technical University in Prague - Faculty of Information Technologies - Software and Web Engineering



Summer semester 2019/2020
Humla v0.3

Overview

- **Introduction**
- Collaborative Filtering
- Content Based Filtering
- Other Recommender Systems
- Other Aspects
- Supplementary Material

Motivation

- Too much information
 - *users are overloaded with information*
 - *many choices available*
- Examples
 - *Thousands of news articles and blog posts every day*
 - *Milions of movies, books, sound tracks*
 - *Many restaurants, hotels, apps, people, ...*
- People can't assess every piece of information!
 - *Need for an intelligent information delivery!*
- Solution
 - **Recommender Systems**
- Relation to Information Retrieval
 - *Many similar approaches on the background*
 - *Goals*
 - IR - *"I know what I am looking for"*
 - RS - *"I am not sure what i am looking for"*

Recommender Systems

- Systems for recommending items to users
 - *books, movies, web pages, songs, ...*
 - *based on users preferences in the past*
- Enhance user experience
 - *assist users in finding information*
 - *reduce search and navigation time*
- Recommendations are on first place on the internet
 - *Google's recommendations*
 - *web pages, advertisements*
 - *recommendations generate 38% more clicks*
 - *Youtube's recommendations*
 - *videos, channels to subscribe*
 - *Netflix's recommendations*
 - *movies*
 - *2/3 watched movies are recommended*
 - *Facebook's recommendations*
 - *friends, posts, advertisements*
 - *Twitter's recommendations*
 - *tweets, people to follow*
 - *Amazon's recommendations*
 - *books, DVDs, electronics, toys, ...*
 - *35% sales from recommendations*

Information Used for Recommendations

- Used information can come from different sources:
 - *browsing and searching - past behavior*
→ e.g. *clicks, logs*
 - *purchase data*
→ e.g. *bought book or movie*
 - *explicit feedback provided by the users*
→ e.g. *likes, rating stars*
 - *comments on products*
→ *requires opinion mining*
 - *demographic data*
 - *context*
 - *relations to other users*
 - *item similarities*
 - ...
- Recommender systems usually exploit information from more sources

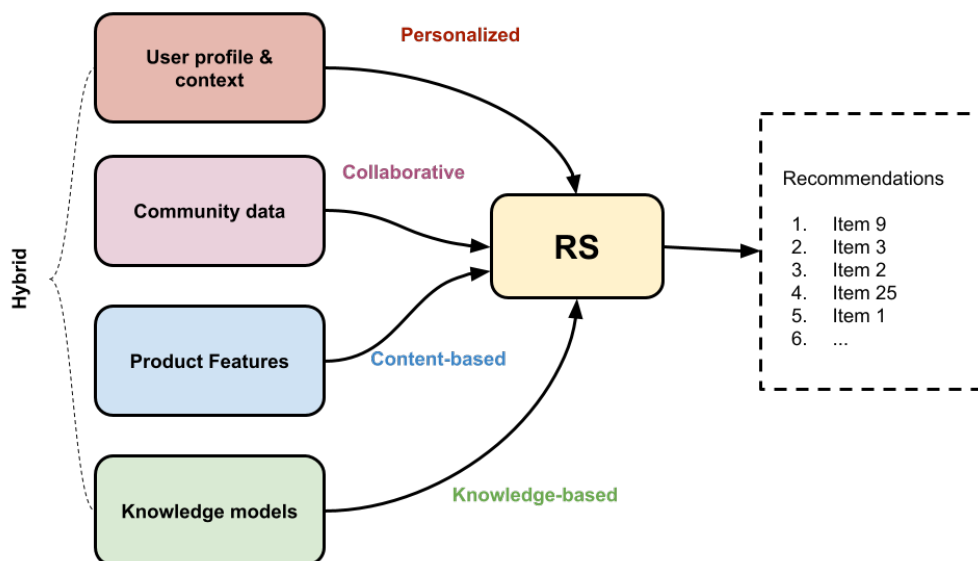
Personalization

- Adapting to the individual needs
- Accommodate the differences between individuals according to their preferences, needs or interests
- Key idea in recommendation systems
- Assumptions
 - *if user A and user B are similar, then in the future user A might be interested in items from user B*
→ *recommend items from similar users*
 - *if user A is interested in item I, then in the future user A might be interested in items similar to item I*
→ *recommend items similar to the items in the basket*

Types of Recommender Systems

- Collaborative filtering
 - *analyze the user-item interactions for each user, and recommend items from similar users*
- Content-based filtering
 - *analyze the content of the items (bought books) and recommend similar content (similar books)*
- Knowledge-based
 - *In situation when limited information is available*
 - *Uses knowledge-base*
 - *which item should be recommended in which context?*
- Other - demographic, ...
- Hybrid
 - *combination of approaches*

Types of Recommender Systems (cont.)



Overview

- Introduction
- **Collaborative Filtering**
- Content Based Filtering
- Other Recommender Systems
- Other Aspects
- Supplementary Material

Collaborative Filtering

- Assumption
 - *User with similar taste in past will have similar taste in future*
- Analyze user-item interactions
 - *likes, watched video, bought book*
 - *do not analyze the content of the items being recommended*
- Method
 - *Produce recommendations by computing the similarity between a user's preferences and preferences of other people*
 - *Suggest new items to users who have similar preferences with others*
- Approaches
 - *User-based - find similar users to me and recommend what they liked*
 - *Item-based - find similar items to those that I have previously liked*
- Widely used in practice

CF Inputs

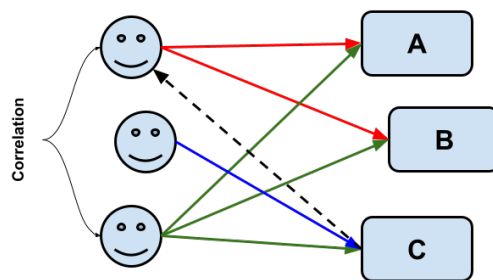
- General lists
 - *List of users*
 - *List of items*
- Associations
 - *Each user has associations to items*
 - *Explicit*
 - *ratings*
 - *Implicit*
 - *purchases*
 - *listen records*
- Metrics/methods
 - *to measure similarity between users*
 - *to select top neighbors*
- Active user

Collaborative Filtering steps

- Steps
 - *Identify relevant input data*
 - *Collect user-item interactions for the users*
 - *ratings (0-5 stars), likes, etc*
 - *Identify set of users most similar to the active user - neighborhood*
 - *Compute user-to-user similarity*
 - *based on their interactions with the items*
 - *Identify items for these similar user*
 - *Generate prediction*
 - *that would be given by the active user*
 - *Recommend items*
- Main challenge
 - *how to compute the similarity between the users*

User-based CF

- Steps
 - For the active user find top similar users
→ neighborhood
 - Detect ratings of neighbors for unrated items
 - Compute average of ratings
 - Recommend items with highest predicted ratings
- In short
 - "You may like it because other liked it"



Example

- Ratings of 5 users for 6 items
- What can we recommend for the user 3?

users\items	1	2	3	4	5	6	mean rating	Cosine(u, u3)	Pearson (u, 3)
u1	7	6	7	4	5	4	5.5	0.956	0.894
u2	6	7		4	3	4	4.8	0.981	0.939
u3		3	3	1	1		2	1.0	1.0
u4	1	2	2	3	3	4	2.5	0.789	-1.0
u5	1		1	2	3	3	2	0.645	-0.817

- Descriptions:
 - For user 3 we can recommend missing items
→ Possibilities - item 1 and 6
 - User 1 and user 2 are the most similar

Similarities

- Cosine similarity
 - Defined as cosine of the angle between the vectors

$$\text{similarity} = \cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|} = \frac{\sum_{i=1}^n A_i \times B_i}{\sqrt{\sum_{i=1}^n (A_i)^2} \times \sqrt{\sum_{i=1}^n (B_i)^2}}$$

- Cosine similarity in action

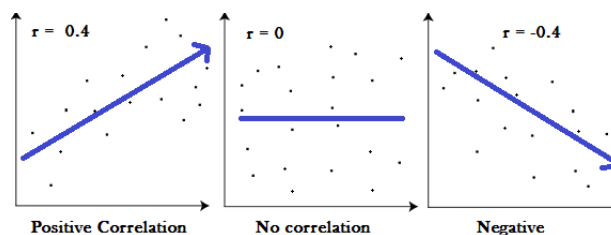
users\items	1	2	3	4	5	6	mean rating	Cosine(u, u3)	Pearson (u, 3)
u1	7	6	7	4	5	4	5.5	0.956	0.894
u3	3	3	1	1			2	1.0	1.0

- cosine(user1, user3)
 - $A \cdot B = 6 \cdot 3 + 7 \cdot 3 + 4 \cdot 1 + 5 \cdot 1$
 - $\|A\| \cdot \|B\| = \sqrt{6^2 + 7^2 + 4^2 + 5^2} \cdot \sqrt{3^2 + 3^2 + 1^2 + 1^2}$
 - $\text{cosine}(\text{user1}, \text{user3}) = A \cdot B / \|A\| \cdot \|B\| = 0.956$

Similarities (cont.)

- Pearson Similarity Metric
 - Computes how correlated are two vectors
 - Similarity in range $[1, -1]$
 - closer to 1 items are correlated, closer to -1 items are negatively correlated
 - negative value means as one variable increases the other decreases
 - positive value means as one variable increases also the other increases
 - in our example, users tends to rank photos with similar ratings

$$\text{corr}(i, j) = \frac{\sum_{i=u}^U (R_{ui} - R_i) (R_{uj} - R_j)}{\sqrt{\sum_{i=u}^U (R_{ui} - R_i)^2} \sqrt{\sum_{i=u}^U (R_{uj} - R_j)^2}}$$



Similarities (cont.)

- Pearson Similarity Metric in Action

users\items	1	2	3	4	5	6	mean rating	Cosine(u, u3)	Pearson (u, 3)
u1	7	6	7	4	5	4	5.5	0.956	0.894
u3	3	3	1	1			2	1.0	1.0

- $\text{corr}(\text{user1}, \text{user3})$

$$- A = (6 - 5.5) * (3 - 2) + (7 - 5.5) * (3 - 2) \dots$$

$$- B = \sqrt{0.5^2 + 1.5^2 + \dots} * \sqrt{1^2 + 1^2 \dots}$$

$$- \text{corr}(\text{user1}, \text{user3}) = A / B = 0.894$$

User-based Recommendation

- How to use the user-user similarity information and predict rating for an item?

- Pearson-weighted average of ratings

– *predicts rating of missing items*

– *predicted ratings are used for recommendation of items*

- Raw rating prediction

$$\text{pred}(a, b) = \frac{\sum_{b \in N^{\text{sim}}(a, b)} r_b}{\sum_{b \in N^{\text{sim}}(a, b)}$$

- Mean-centered prediction

$$\text{pred}(a, b) = r_a + \frac{\sum_{b \in N^{\text{sim}}(a, b)} (r_{bp} - r_b)}{\sum_{b \in N^{\text{sim}}(a, b)}$$

User-based Recommendation (cont.)

- What can we recommend for the user 3?

users\items	1	2	3	4	5	6	mean rating	Cosine(u, u3)	Pearson (u, 3)
u1	7	6	7	4	5	4	5.5	0.956	0.894
u2	6	7		4	3	4	4.8	0.981	0.939
u3		3	3	1	1		2	1.0	1.0
u4	1	2	2	3	3	4	2.5	0.789	-1.0
u5	1		1	2	3	3	2	0.645	-0.817

- In action
 - by selecting top-2 similar users

$$\rightarrow \text{user3, item1} = \frac{7 * 0.894 + 6 * 0.939}{0.894 + 0.939} = 6.49$$
 - $$\rightarrow \text{user3, item6} = \frac{4 * 0.894 + 4 * 0.939}{0.894 + 0.939} = 4$$

Item-based CF

- Computes similarity between items (columns in the rating matrix)
 - Use this similarity to predict ratings
 - More computationally efficient
 - number of items \ll number of user
 - Similar similarity metrics can be used
 - Prediction is computed in the same way as in the user-based

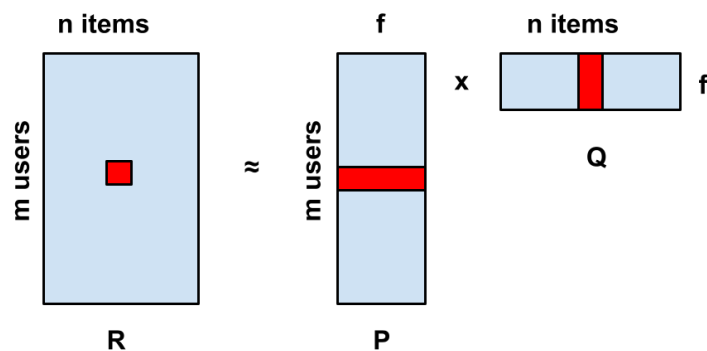
users\items	1	2	3	4	5	6
u1	7	6	7	4	5	4
u2	6	7		4	3	4
u3		3	3	1	1	
u4	1	2	2	3	3	4
u5	1		1	2	3	3
sim(1, i)	1	0.735	0.912	-0.848	-0.813	-0.990
sim(6, i)	-0.990	-0.622	-0.912	0.829	0.730	1

Matrix Factorization CF

- Model-based recommendation
- Main idea
 - Discover latent factors of users or items
 - Use those latent factors for recommendations
 - Assumption: Ratings can be inferred from a model put together from a smaller number of parameters
- Latent factors
 - features that describe recommended items or users
 - e.g. categories of movies
 - can be discovered automatically
- Approaches
 - Singular Value Decomposition (SVD)
 - Matrix Factorization

Matrix Factorization CF (cont.)

- Basic Matrix Factorization



- Description
 - R - rating matrix (m users, n items)
 - P - user features matrix (m users, f features/latent factors)
 - W - item features matrix (n items, f features/latent factors)
 - Rating can be computed as
 - $\rightarrow r_{ui} \approx q_i^T \times p_u$

Matrix Factorization CF - Example

- Rating matrix

users\items	item1	item2	item3
u1	5	3	4
u2	?	2	4
u3	4	2	?

- P

users/features	f1	f2	f3	f4	f5
u1	0.276	-0.377	-1.262	-1.548	0.473
u2	0.399	-0.527	-0.289	-1.516	0.737
u3	0.225	-0.291	-1.067	-1.225	0.374

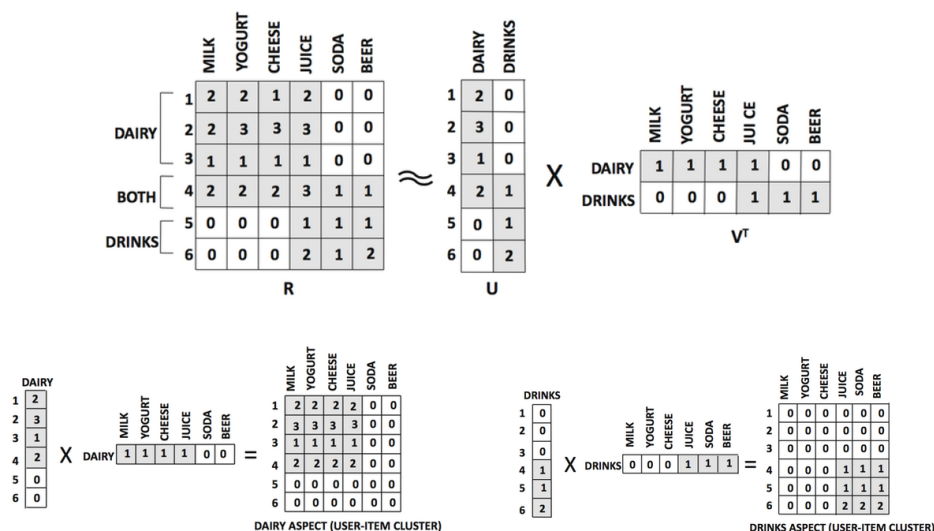
- W

items/features	item1	item2	item3
f1	0.302	0.143	0.434
f2	-0.405	-0.202	-0.575
f3	-1.399	-0.923	-0.468
f4	-1.678	-0.940	-1.720
f5	0.512	0.239	0.796

- Rating prediction: $r(\text{user3}, \text{item3}) = P_3 f \times W_3 \beta = 3.168$

Matrix Factorization CF - Example 2

- Product Ratings



Matrix Factorization CF - Example 3

- *Movie Ratings*

User ID	Top Gun	Rambo	Young Guns	Lonesome Dove
1	5	4	1	1
2	4	5	1	1
3	1	1	5	5
4	1	1	4	5
5	1	1	5	4

- *Feature Vectors for movies*

	Feature 1	Feature 2
	Military	Western
Top Gun	3.836413	0.5803331
Rambo	7.744796	0.3196278
Young Guns	1.217036	7.829875
Lonesome Dove	1.120779	9.8565764

* "Nonnegative Matrix Factorization and Recommender Systems." R-bloggers. N.p., 24 Oct. 2012. Web. 28 Apr. 2017.

Problems with Collaborative Filtering

- Cold start
 - *needs to have enough other users already in the system to find match*
- Sparsity
 - *if there are many items to be recommended, even if there are many users, the user/ratings matrix is sparse, and it is hard to find users that have rated the same items*
- First rater
 - *new not rated items are not recommended*
- Popularity bias
 - *can not recommend items to someone with unique tastes*
 - *specific items are always popular*
- Lack of Transparency
 - *recommendations can not be easily explained, why an item was recommended*

Overview

- Introduction
- Collaborative Filtering
- **Content Based Filtering**
- Other Recommender Systems
- Other Aspects
- Supplementary Material

Content Based Filtering

- Produce recommendations by analyzing the content (metadata) of the items
 - *metadata - categories, tags, textual content, authors, etc.*
 - *automatically extracted fetures - e.g. text mining*
- Approach
 - *Build representation for each item*
 - *Represent items usually as a vector*
 - *Content-based learning of user profiles*
 - *Can be represented in many ways*
 - *similar vector as for the items*
 - *set of keywords*
 - *classification or regression models*
 - *Filtering and recommendations*
 - *Identify candidates*
 - *Measure the similarity between the user profile vector and each item vector*
 - *Classify items using user profiles as classifiers*
 - *Recommend top-N most relevant items to the user*

Content Based Filtering - Example

- Items

	Action	Sci-fi	Comedy	Romance	Sport	Drama
item1	1	1	0	0	0	0
item2	1	0	1	0	0	0
item3	0	0	1	1	0	0
item4	0	0	0	0	1	1
item5	0	0	1	0	1	0
item6	0	1	0	0	0	1
item7	1	1	1	0	0	0

- User profile

	Action	Sci-fi	Comedy	Romance	Sport	Drama
user1	1	0	0	0	0	1

- Similarity measures
 - *item1, item2, item4, item6, ...*

Content Based Filtering - Example 2

- Rule-based classifier

song\cat.	Drums	Guitar	Beat	Classical	Symphony	Orchestra	L/D
1	1	1	1	0	0	0	D
2	1	1	0	0	0	1	D
3	0	1	1	0	0	0	D
4	0	0	0	1	1	1	L
5	0	1	0	1	0	1	L
6	0	0	0	1	1	0	L
t1	0	0	0	1	0	0	?
t2	1	0	1	0	0	0	?

- Rules (supp=0.3, conf=0.75)
 - Rule 1: {Classical} -> Like (50%, 100%)
 - Rule 2: {Symphony} -> Like (33%, 100%)
 - Rule 3: {Classical, Symphony} -> Like (33%, 100%)
 - Rule 4: {Drums, Guitar} -> Dislike (33%, 100%)
 - Rule 5: {Drums} -> Dislike (33%, 100%)
 - Rule 6: {Beat} -> Dislike (33%, 100%)
 - Rule 7: {Guitar} -> Dislike (50%, 75%)

- Classifications

- t1 – Rule 1 – Like
- t2 – Rule 5, Rule 6 – Dislike

Advantages of Content Based Filtering

- No need for data from other users
 - *No cold start or sparsity problems*
- Tastes of unique users are met
- Can recommend new and unpopular items
 - *First rater problem is solved*
- Can provide explanation why an item was recommended
 - *by listing content features that caused an item to be recommended*

Disadvantages of Content Based Filtering

- Content should be encoded as meaningful features
- Users tastes must be possible to represent as vectors
- Problems with new users with no profile
- Does not exploit the collective intelligence information
 - *how other users interacted with the item*
- Difficult to implement serendipity
- Easy to overfit

Overview

- Introduction
- Collaborative Filtering
- Content Based Filtering
- **Other Recommender Systems**
- Other Aspects
- Supplementary Material

Knowledge-Based RS

- Applications
 - *expensive items, not frequently purchased, few ratings*
 - *time span important (e.g., technological products)*
 - *ratings may be time-sensitive*
 - *"The ratings on an old car or computer are not very useful for recommendations because they evolve with changing product availability"*
 - *explicit requirements of user*
 - *interactivity is a crucial component of such systems*
- Limitations of other approaches
 - *collaborative filtering unusable – not enough data*
 - *content based – similarity not sufficient*
- Requires functional/domain knowledge
 - *about how a particular item meets a particular user need*

Knowledge-Based RS (cont.)

- Types
 - *Constraint-based recommender systems*
 - users typically specify requirements or constraints (e.g., lower or upper limits) on the item attributes
 - domain-specific rules are used to match the user requirements
 - e.g. "Cars before year 1970 do not have cruise control."
 - e.g. "Older investors do not invest in ultrahigh-risk products."
 - process is interactively repeated until the user arrives at her desired results
 - *Case-based recommender systems*
 - specific cases are specified by the user as targets
 - similarity metrics are defined on the item attributes to retrieve similar items to these targets
 - returned results are often used as new target cases
 - interactive process is used to guide the user towards the final recommendation
- Interactions
 - *Conversational systems*
 - user preferences are determined in the context of a feedback loop
 - *Search-based systems*
 - a preset sequence of questions
 - *Navigation-based recommendation*

Constraint-Based RS - Example

- Buying homes

Item	Beds	Baths	Locality	Type	Floor Area	Price
1	3	2	Bronx	Townhouse	1600	220,000
2	5	2.5	Chappaqua	Split-level	3600	973,000
3	4	2	Yorktown	Ranch	2600	630,000
4	2	1.5	Yorktown	Condo	1500	220,000
5	4	2	Ossining	Colonial	2700	430,000

- Customer-specified attributes
 - *Marital-status* (categorical), *Family-Size* (numerical), *suburban-or-rural* (binary), *Min-Bedrooms* (numerical), *Max-Bedrooms* (numerical), *Max-Price* (numerical)
- Domain knowledge for mapping of requirement into the attributes
 - *Suburban-or-rural*=Suburban -> *Locality*= {List of relevant localities}
 - *Marital-status*=single -> *Min-Bedrooms*≤5
 - *Family-Size*≥5 -> *Min-Bedrooms*≤3
 - *Min-Bedrooms*≥3 -> *Price*≥100,000

Other RS

- Context-sensitive RS
 - *context - additional information that defines the specific situation under which recommendations are made*
 - Time - morning, evening, weekdays, weekends, holidays
 - Location - recommendation for a restaurant in his locality
 - Social information - e.g. social circles can affect the recommendation process
- Ensemble or Hybrid RS
 - *Weighted*
 - scores of several recommender systems are combined into a single unified score
 - *Switching*
 - switches between various recommender systems depending on current needs
 - *Cascade*
 - one recommender system refines the recommendations given by another
 - *Feature augmentation*
 - the output of one recommender system is used to create input features for the next
 - *Mixed*

Overview

- Introduction
- Collaborative Filtering
- Content Based Filtering
- Other Recommender Systems
- Other Aspects
- Supplementary Material

Evaluation of RS

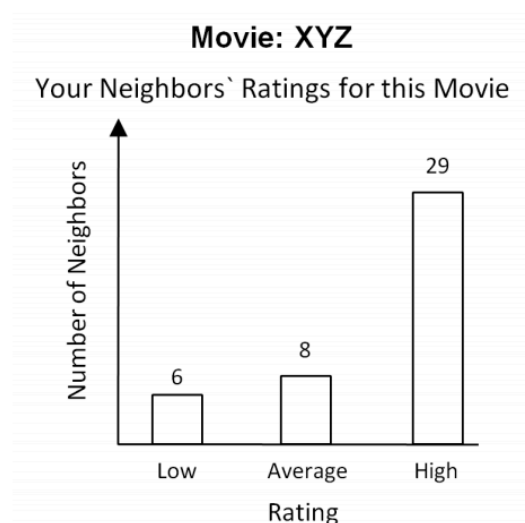
- Evaluations - difficult to properly measure the quality
 - many approaches
 - many metrics
- Approaches
 - off-line
 - simulations of users and their behavior
 - issues with overfitting
 - the goal is to filter out inappropriate approaches
 - user-studies
 - small group of real users
 - expensive
 - on-line
 - AB testing
 - can be risky - user dissatisfaction
- Metrics
 - Root Mean Squared Error (RMSE) - normalized and averaged
 - Precision, Recall (at N)
 - F-measure, Area Under the ROC Curve (AUC)
 - Normalized Cumulative Discounted Gain (NDCG)
 - Click-through rate (CTR)
 - Coverage - Novelty, Serendipity, Diversity

Explanations

- Situation
 - recommendations: selection (ranked list) of items
 - explanations: (some) reasons for the choice
- Why explanations?
 - transparency, trustworthiness, validity, satisfaction (users are more likely to use the system)
 - persuasiveness (users are more likely to follow recommendations)
 - effectiveness, efficiency (users can make better/faster decisions)
 - education (users understand better the behaviour of the system, may use it in better ways)
- Examples
 - knowledge-based recommender
 - "Because you, as a customer, told us that simple handling of car is important to you, we included a special sensor system in our offer that will help you park your car easily."
 - recommendations based on item-similarity
 - "Because you watched X we recommend Y"

Explanations - Example

- Collaborative filtering



Movie: XYZ
Personalized Prediction: ****
Your Neighbors' Ratings for this Movie

Rating	Number of Neighbors
★	2
★★	4
★★★	8
★★★★	20
★★★★★	9

* Explaining collaborative filtering recommendations, Proceedings of the 2000 ACM Conference on Computer Supported Cooperative Work (CSCW)

Attacks

- Most vulnerable
 - Collaborative filtering RS
- Reasons for attacks
 - make the system worse (unusable)
 - influence rating (recommendations) of a particular item
 - push attacks – improve rating of "my" items
 - nuke attacks – decrease rating of "opponent's" items
- Approach
 - creating a set of fake feedbacks from many different users
 - shilling attack - users become shills in the attack

Attacks (cont.)

- Example

	Items						
	1	2	3	4	5	6	7
<i>a</i>	+	-		+	+		+
<i>b</i>	-	+	+	-	-		-
<i>c</i>	+	-	+		-	-	-
<i>d</i>	-	+	+	-			
<i>e</i>	-		-	-	-		-
<i>f</i>	+	-	+	+	+		+
<i>g</i>		-	+	+	-	-	+
<i>h</i>	+	-	+	+	+		?
<i>i</i>	+	-	+		-	-	-
<i>j</i>	-	+	+	-			-
<i>k</i>	-		-	-	-		-
<i>l</i>	+	-	+	+	+		-
<i>m</i>		-	+	+	-	-	-

Users

Authentic profiles

Target profile

Attack profiles

Attacks - Types

- Levels
 - more knowledge about system -> more efficient attack
- Types
 - random attack
 - generate profiles with random values
 - requires global mean
 - average attack
 - effective attack on memory-based systems (average ratings -> many neighbors)
 - requires global mean and mean of each item
 - bandwagon attack
 - high rating for "blockbusters", random values for others
 - requires to know which items are the most popular
 - segment attack
 - insert ratings only for items from specific segment
 - attack item-based collaborative filtering
 - love/hate attack
 - the nuked item is set to the minimum rating value, whereas the other items are set to the maximum
 - reverse bandwagon
 - widely disliked items are used as filler items to mount the attack
 - probe attack
 - a seed profile is created by the attacker, and the predictions generated by the recommender system are used to learn related items and their ratings

Attacks - Detection

- Attack detections
 - *Unsupervised attack detection algorithms*
 - *ad hoc rules are used to detect fake profiles*
 - *to identify the key characteristics of attack profiles that are not similar to genuine profiles*
 - *if a profile (or significant portion of it) is identical to many other profiles*
 - *Number of prediction differences, Degree of disagreement with other users, Rating deviation from mean agreement, Degree of similarity with top-k neighbors*
 - *Supervised attack detection algorithms*
 - *classification models to detect attacks*
 - *number of profiles to which a given user profile is identical can be used as a feature for that user profile*
- Robust recommender systems
 - *preventing automated attacks with CAPTCHAs*
 - *detections using metrics, clusterings, neighborhoods, more complex algorithms etc.*

Overview

- Introduction
- Collaborative Filtering
- Content Based Filtering
- Other Recommender Systems
- Other Aspects
- **Supplementary Material**

Euclidean Based Proximity Measure

- Represent query and document in a vector space model
 - represent a query as a vector
 - represent a document as a vector
- Compute the distance of the two vectors in a multi-dimensional space
- An Euclidean distance of two vectors

$$d(\mathbf{p}, \mathbf{q}) = d(\mathbf{q}, \mathbf{p}) = \sqrt{(q_1 - p_1)^2 + (q_2 - p_2)^2 + \dots + (q_n - p_n)^2} = \sqrt{\sum_{i=1}^n (q_i - p_i)^2}$$

- Euclidean distance is a bad idea
 - if a query term occurs more times in the document, the distance will be larger
 - distance is large, although the distribution of terms are similar

Cosine Based Proximity Measure

- Use angle instead of distance
 - compute cosine of the angle between the query and document vectors
- Cosine similarity of two vectors (query vs. document)

$$\text{similarity} = \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i \times B_i}{\sqrt{\sum_{i=1}^n (A_i)^2} \times \sqrt{\sum_{i=1}^n (B_i)^2}}$$

- Long and short vectors now have comparable weights
- If two documents are identical, then the angle is 0
 - cosine of 0 angle = 1
 - similarity of two identical vectors is 1
- Rank the documents according to the angle of the query

Precision and Recall

- Precision
 - *fraction of retrieved documents that are relevant*
 - $P(\text{relevant}|\text{retrieved})$
- Recall
 - *fraction of relevant documents that are retrieved*
 - $P(\text{retrieved}|\text{relevant})$

Contingency table:

I	Relevant	Non-relevant
Retrieved	tp (true positive)	fp (false positive)
Not retrieved	fn (false negative)	tn (true negative)

Precision and Recall

I	Relevant	Non-relevant
Retrieved	tp (true positive)	fp (false positive)
Not retrieved	fn (false negative)	tn (true negative)

- tp - true positive
 - *relevant documents that are retrieved*
- fp - false positive
 - *non-relevant documents that are retrieved*
- fn - false negative
 - *non-retrieved relevant documents*
- tn - true negative
 - *non-retrieved non-relevant documents*

Precision and Recall Measure

/	Relevant	Non-relevant
Retrieved	tp (true positive)	fp (false positive)
Not retrieved	fn (false negative)	tn (true negative)

- Precision

$$- P = \frac{tp}{(tp + fp)}$$

- Recall

$$- R = \frac{tp}{(tp + fn)}$$

Accuracy

- Assume

- an engine for a given query classifies each document as "relevant" or "non relevant"

- The accuracy

- the fraction of these classifications that are correct

$$- \frac{(tp + tn)}{(tp + fp + fn + tn)}$$

- Extremely skewed

- normally over 99.9% of the documents are in the non-relevant category

- Widly used measure in evaluating machine learning systems

F-Measure

- Also known as F-score or F1 score
- Assesses precision/recall tradeoff

$$F = 2 \times \frac{(Precision \times Recall)}{(Precision + Recall)}$$

