

Web Data Mining

Lecture 3: Data Access and Acquisition Methods 2

Jaroslav Kuchař & Milan Dojčinovski

jaroslav.kuchar@fit.cvut.cz, milan.dojchinovski@fit.cvut.cz



Czech Technical University in Prague - Faculty of Information Technologies - Software and Web Engineering



Summer semester 2019/2020
Humla v0.3

Overview

- Parsing and Extracting
- Web Scraping
- Accessing Specific Content

Data Types

- Unstructured
 - *Lacks any structure*
 - *e.g. free texts, documents, multimedia, ...*
- Semistructured
 - *Inconsistent structure*
 - *Often self-describing (e.g. key-value pairs), flexibility*
 - *e.g. CSV, XML, ...*
- Structured
 - *Conform to the predefined model*
 - *Predefined schema*
 - *Highly structured*
 - *e.g. numbers, dates, structured entities, ...*

Recall: Parsing

- Parsing the content of the HTTP payload
 - *extracting content for indexing*
 - *extracting links to be added to the frontier*
 - *extracting additional crawling and indexing directives*
 - *headers Cache-Control, Content-Type, X-Robots-Tag, ...*
- HTML code very often contains invalid markup
 - *unclosed elements, unencoded special characters, missing required attributes, improperly nested tags, missing quotes, ...*
- Bad HTML markup should be fixed
 - *a preprocessing step is required to clean up the HTML*
 - *many tools available, tidy - a tool provided by W3C*

HTML markup fix

- Tidy:

```
1 from tidylib import tidy_document
2 document, errors = tidy_document(''<p>f&otilde;o '',
3 options={'numeric-entities':1})
4 print(document)
5 print(errors)
```

- Output:

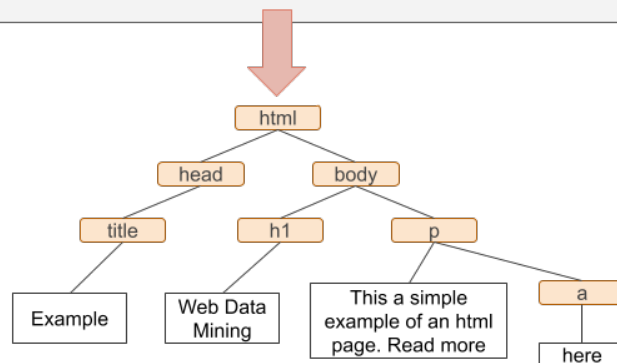
```
1 <!DOCTYPE html PUBLIC "-//W3C//DTD HTML 4.01//EN">
2 <html>
3   <head>
4     <title></title>
5   </head>
6   <body>
7     <p>
8       f&otilde;o 
9     </p>
10  </body>
11 </html>
```

```
1 line 1 column 1 - Warning: missing <!DOCTYPE> declaration
2 line 1 column 1 - Warning: inserting missing 'title' element
3 line 1 column 18 - Warning: <img> lacks "alt" attribute
```

Illustration of the DOM tree model

- Parsing of a simple HTML page

```
<html>
<head>
<title>Example</title>
</head>
<body>
<h1>Web Data Mining course</h1>
<p>This a simple example of an html page. Read more
<a
href="http://www.w3schools.com/html/html_examples.asp">here</a>
</p>
</body>
</html>
```



Extraction of Relevant Information

- Web Information Extraction
 - *Natural language text processing*
 - *Extracting structured data from Web pages*
- Basic approaches
 - *Text fields identification*
 - *Title, meta information, text blocks*
 - *Each can have different importance (e.g. HTML title, h1, h3, ...)*
 - *Anchor's texts extraction*
 - *They are acting as a short description of the target.*
 - *Important for search engines.*
- Detection of relevant structures/information
 - *Recognizers*
 - *Wrappers*
 - *Automatic approaches*
- Other approaches
 - *Machine readable annotations, JS variables, API calls, ...*

Recognizers

- Follow a procedure to find a piece of information based on its appearance
 - *e.g. e-mail addresses, phone numbers or street addresses*
- Most of nowadays crawlers and search engines are able to automatically collect such information.
- Can be executed as a set of regular expression patterns:

```
1 import re
2 phonePattern = re.compile(r'^(\d{3})-(\d{3})-(\d{4})$')
3 phonePattern.search('800-555-1212').groups()
4 # ('800', '555', '1212')
```




```
1 import re
2 emailPattern = re.compile(r'([a-zA-Z0-9._%+-]+)@([a-zA-Z0-9.-]+\.[a-zA-Z]{2,})')
3 emailPattern.search('user@sub.example.com').groups()
4 # ('user', 'sub.example.com')
```

Wrappers

- Allows to a semi-structured web data source be consulted as if it was a common database.
- Approaches
 - *Manual*
 - *Observing the source code, finding patterns and writing the crawler*
 - *Not scalable*
 - *Induction*
 - *Supervised learning approach, semiautomatic*
 - *Set of manually labeled pages*
 - *Automatic*
 - *Unsupervised approach*
 - *Automatically finds patterns (e.g. repetitive structures, visual aspects, ...)*
 - *Can scale*

Wrapper Induction Approaches

- Learning from a set of labeled instances
 - *Manual annotations of informations that are required for extraction.*
 - *e.g. product name, price and associated attributes*
- Approach
 - *Tree based representations of annotated information*
 - *Possibilities for generalizations/pruning*
 - *Learning of extraction rules:*
 - *e.g. `SkipTo("")`*



- Limitations
 - *Manual annotation is not possible in large scale.*
 - *Web is dynamic. Cost of maintenance.*

Automatic Extraction

- Based on mining repeated patterns in multiple structures.
- Uses matching approaches
 - *String matching*
→ e.g. edit distances - minimum number of changes to change one string to another one ("Page 1" vs "Page 2", "Notebook" vs "CPU")
 - *Tree matching*
→ minimum set of operation to transform trees
- Building and processing trees
 - *Pure DOM based tree and tree pruning*
 - *Visual aspects can influence the importance of parts*
→ e.g. size, colors, visibility, ...
- Patterns
 - *Repetitive tree structures with the same content can represent headers, footers, menus or ads*
 - *Varying subtrees across multiple pages identify main content*
→ *High density of textual nodes can represent the main textual content*
 - *Structures close to each other but different colors represents different things*
 - ...
- Tools
 - *boilerpipe, readability,*
 - *diffbot* - analyze the visual layout, *dragnet* - rule-based models, *octoparse*
 - *mercury, fathom*

Overview


- Parsing and Extracting
- **Web Scraping**
- Accessing Specific Content

Web Scraping

- Web Scraping vs Crawling
 - Web Crawling
 - usually provides data for indexing, search engines etc.
 - following links, finding new pages and collecting content
 - Web Scraping
 - for pages without API
 - extracting structured information from a web page
 - usually specific for the target website and structured data
 - more focused process
 - Many overlapping actions
 - execute JavaScript, emulate human user behavior, submit forms, log in to a website etc.

Web Scraping (cont.)

- Main process
 1. Pick a URL
 - e.g. one from the predefined or collected list of URLs
 2. Collect the URL content
 - HTTP request/response
 3. Parse and extract desired information from the structure
 - using predefined rules or automatic detection mechanisms



Web Scraping Example - Fetching Data

- Sending HTTP requests

```
1 import requests
2
3 url="https://en.wikipedia.org/wiki/Web_scraping"
4
5 # Make a GET request to fetch the raw HTML content
6 html_content = requests.get(url).text
7
8 print(html_content)
```

- Output:

```
1 <!DOCTYPE html>
2 <html class="client-nojs" lang="en" dir="ltr">
3 <head>
4 <meta charset="UTF-8"/>
5 <title>Web scraping - Wikipedia</title>
6 ...
```

Parsing Data - BeautifulSoup

- BeautifulSoup

- Python library for pulling data out of HTML and XML files

- Parsers

- lxml

- fast

- parses broken HTML

- external dependency

- html.parser

- internal dependency in Python

- html5lib

- parses broken HTML

- external dependency

- Parsing

- navigating HTML tree

- children, next_siblings, previous_sibling, parent

- selecting elements

- find, find_all

- using CSS selectors, ...

Web Scraping Example - Parsing Data

```
1 | # Parse the html content
2 | soup = BeautifulSoup(html_content, "lxml")
3 | # soup = BeautifulSoup(html_content, "html.parser")
4 | print(soup.prettify()) # print the parsed data of html
5 |
6 | # accessing DOM tree elements
7 | print(soup.title)
8 | print(soup.title.text)
9 |
10 | # finding elements
11 | print(soup.find_all('section', class_='categories'))
12 | print(soup.find(id='product-price').children)
13 |
14 | # using CSS selectors
15 | print([e.text for e in soup.select("h2:has(> span#See_also) + div > ul > li

1 | ....
2 |
3 | ['Archive.is', 'Comparison of feed aggregators', 'Data scraping',
4 | 'Data wrangling', 'Importer', 'Job wrapping', 'Knowledge extraction',
5 | 'OpenSocial', 'Scraper site', 'Fake news website', 'Blog scraping',
6 | 'Spamdexing', 'Domain name drop list', 'Text corpus',
7 | 'Web archiving', 'Blog network', 'Search Engine Scraping',
8 | 'Web crawlers']
```

Web Scraping Example - Parsing Data 2

```
1 | import re
2 | # Parse the html content
3 | soup = BeautifulSoup(html_content, "lxml")
4 | # soup = BeautifulSoup(html_content, "html.parser")
5 |
6 | soup.find_all('a', {'href':re.compile('*.html') }).attrs['href']
7 |
8 | soup.find_all(lambda t: len(t.attrs)==3)

1 | ....
2 |
3 | [ ... ]
```

Overview

- Parsing and Extracting
- Web Scraping
- Accessing Specific Content

Crawling Deep Web

- Deep Web
 - The content hidden for crawling engines - e.g. behind HTML forms.
 - Usually user databases, registration-required web forums, web mail pages, online banking, pages behind paywalls (video on demand), ...

<https://www.cisoplatform.com/profiles/blogs/surface-web-deep-web-and-dark-web-are-they-different>



Crawling Deep Web (cont.)

- Categories
 - *Contextual web* - content dependent on the context (e.g. location)
 - *Dynamic pages* - response to a query (e.g. search)
 - *Restricted access* - technical restrictions (e.g. robots.txt, CAPTCHA)
 - *Non-HTML content* - multimedia or other file formats
 - *Private Web* - login required pages
 - *Scripted content* - AJAX, Flash, ...
 - *Unlinked* - no backlinks
 - *Hidden* - using only specific software, protocols (e.g. Tor)

Crawling Dynamic pages

- Detection of forms
 - e.g. search form
- ```
1 | <input type="search">
```
- Query generation
    - *pre-computing submissions for each HTML form*
      - *Using known entities from a knowledge base or search engine query logs*
  - URL generation
    - *replacing parameters in URLs*
      - e.g. <https://www.google.com/search?q=deep+web>
  - Empty page filtering
    - *for detection of irrelevant generations*
  - Deduplications of results

## Selenium

- Selenium
  - *primarily for automating web applications for testing purposes*
  - *can be used for scraping/crawling as well*

```
1 # Example for forms
2 from selenium.webdriver import Firefox
3 from selenium.webdriver.firefox.options import Options
4 opts = Options()
5 opts.set_headless()
6 browser = Firefox(options=opts)
7 browser.get('https://duckduckgo.com')
8 search_form = browser.find_element_by_id('search_form_input_homepage')
9 search_form.send_keys('python')
10 search_form.submit()
11 results = browser.find_elements_by_class_name('result__a')
12 print(results[0].text)
13 browser.close()
```

## Crawling Restricted/Private pages

- Limited possibility to access the data
  - *Login forms with email confirmations*
  - *CAPTCHA, reCAPTCHA, ...*
    - *Completely Automated Public Turing test to tell Computers and Humans Apart (CAPTCHA)*
    - *Visuals, sounds, ...*
    - *Secondary goals - image annotations, digitizing texts, ...*
- Approaches
  - *Special agents for session/cookie based approaches*
    - *Automatic or manual registrations/logins*
  - *CAPTCHA solvers, ...*
  - *Simulations of real human behavior, ...*

## Web Scraping Example - "Real Browser"

- Sending HTTP requests

```
1 import requests
2
3 headers = {'user-agent': 'Mozilla/5.0 (Macintosh; Intel Mac OS X 10_1
4 url="https://en.wikipedia.org/wiki/Web_scraping"
5
6 # Make a GET request to fetch the raw HTML content
7 html_content = requests.get(url, headers=headers).text
8
9 print(html_content)
```

- Output:

```
1 <!DOCTYPE html>
2 <html class="client-nojs" lang="en" dir="ltr">
3 <head>
4 <meta charset="UTF-8"/>
5 <title>Web scraping - Wikipedia</title>
6 ...
```

## Web Scraping Example - Session

- Sending HTTP requests

```
1 import requests
2
3 # requests.get(url, cookies={...})
4
5 session = requests.Session()
6 login_data = {'user': 'user', 'password': 'password'}
7 session.post('https://example.com/login', login_data)
8
9 result = s.get('https://example.com/data')
10
11 print(result.text)
```

- Output:

```
1 <!DOCTYPE html>
2 <html class="client-nojs" lang="en" dir="ltr">
3 <head>
4 <meta charset="UTF-8"/>
5 <title>Web scraping - Wikipedia</title>
6 ...
```

## Crawling AJAX applications

- Modern applications produce content dynamically, faster and richer

```
1 <div id="main">
2 </div>
3 <script>
4 for(var i=0; i<=5; i++){
5 var data = "Link"+i+"
";
6 document.getElementById('main').innerHTML += data;
7 }
8 </script>
```

Link0

Link1

Link2

Link3

Link4

Link5

## Crawling AJAX applications

- AJAX application's content is produced dynamically by the browser
  - dynamically fetch data from the server using JavaScript and display to the user
  - e.g., by clicking on button or link
  - this dynamically created content is not visible to crawlers
- Typical scenario
  - user opens a web page, e.g., <http://example.com/>
  - user clicks on link products
  - a JavaScript code is executed, which dynamically changes the URL to a pretty URL <http://example.com/#products> and content of the current page
- Problem: a browser can execute JavaScript and produce content on-the-fly, the crawler generally can have limitations
- Headless browsers solutions
  - No GUI, emulated DOM, ...
  - can be expensive, resource-wise, ...

## Crawling AJAX applications

- Another problem: only the URL part before the hashtag (#) is processed from the server
  - hash fragments are not part of the HTTP requests
  - hash fragments are not sent to the server
- HTTP request containing hash fragment

```
1 | curl -v https://mail.google.com/mail/u/0/#inbox
2 |
3 | > GET /mail/u/0/ HTTP/1.1
4 | > Host: mail.google.com
5 | > User-Agent: curl/7.30.0
6 | > Accept: */*
7 |
8 | < ...HTTP response...
```

- Note the path `/mail/u/0/` in the HTTP request
  - it does not contain the hash fragment

## Making AJAX Applications Crawlable

- Solution:
  - The crawler detects a pretty URL
    - a URL containing hash fragment beginning with !
    - e.g., <https://mail.google.com/mail/u/0/#!inbox>
  - The crawler transforms the pretty URL to a ugly URL
    - the `#!` is replaced with `_escaped_fragment_`
    - the `_escaped_fragment_` becomes part of the query parameters
    - e.g., <https://mail.google.com/mail/u/0/#!inbox> becomes [https://mail.google.com/mail/u/0/?\\_escaped\\_fragment\\_=inbox](https://mail.google.com/mail/u/0/?_escaped_fragment_=inbox)
  - The crawler requests the ugly URL
  - The server interprets the ugly URL and serves the content
  - The crawler processes the content
- Reference
  - see [Google AJAX Crawling](#) for more details (officially deprecated as of October 2015)

## Spider Traps

- Spider Trap - a set of web pages that create an infinite number of URLs for crawling
  - *finding new previously unvisited URLs for infinite amount of time*
- Harmful for crawlers
  - *waste of bandwidth and storage to download and store duplicate and useless data*
  - *waste of server's bandwidth and space with fake information*
- Creation of spider's traps for good causes
  - *to catch spambots, which send emails on behalf of others*
  - *to catch crawlers, which waste a website's bandwidth*
- Example - Calendar with "next day/month/week/year" link
  - *creation of dynamic pages with links, which point to next year or month*
  - *<http://cal.org/01/2014/> has "next" link pointing to <http://cal.org/01/2015/>*
  - *again <http://cal.org/01/2015/> has "next" link pointing to <http://cal.org/01/2016/>*

## Honeypot Traps

- Honeypot Traps
  - *In the form of links which are not visible to the typical user on the browser*
    - *e.g. setting the CSS as **display: none***
  - *hidden form inputs*
    - *e.g. predefined random/empty value expected (e.g. [GitHub login page](#))*
    - *specific name, that the crawler fill in e.g. **username**, **password***
- Crawler
  - *can try to scrape the information from the link*
  - *can be detected and owner blocks the source IP address*
- Detection is not easy



## Web Scraping Example - Proxies

- Sending HTTP requests  
– <https://free-proxy-list.net/>

```
1 import requests
2 url = 'https://httpbin.org/ip'
3 proxies = {
4 "http": 'http://62.244.49.202:52323',
5 "https": 'http://62.244.49.202:52323'
6 }
7 response = requests.get(url, proxies=proxies)
8 print(response.json())
```

- Output:

```
1 {'origin': '62.244.49.202'}
```

## Other Tools, Demos, Examples

- Headless browsers tools  
– <https://github.com/GoogleChrome/puppeteer>  
– <https://github.com/emadehsan/thal>  
– <https://nickjs.org/>

```
1 const nick = new Nick()
2 ;(async () => {
3 const tab = await nick.newTab()
4 await tab.open("news.ycombinator.com")
5 await tab.untilVisible("#hnmain")
6 await tab.inject("../injectables/jquery-3.0.0.min.js")
7 const hackerNewsLinks = await tab.evaluate((arg, callback) => {
8 const data = []
9 $(".athing").each((index, element) => {
10 data.push({
11 title: $(element).find(".storylink").text(),
12 url: $(element).find(".storylink").attr("href")
13 })
14 })
15 callback(null, data)
16 })
17 console.log(JSON.stringify(hackerNewsLinks, null, 2))
18 })()
19 .then(() => {
20 console.log("Job done!")
21 nick.exit()
22 })
23 .catch((err) => {
24 console.log(`Something went wrong: ${err}`)
25 nick.exit(1)
26 })
```

## Other Tools, Demos, Examples

- <https://blog.phantombuster.com/web-scraping-in-2017-headless-chrome-tips-tricks-4d6521d695e8>

```
1 if (await tab.isVisible(".captchaImage")) {
2 // Get the URL of the generated CAPTCHA image
3 // Note that we could also get its base64-encoded value and solve it too
4 const captchaImageLink = await tab.evaluate((arg, callback) => {
5 callback(null, $(".captchaImage").attr("src"))
6 })
7
8 // Make a call to a CAPTCHA solving service
9 const captchaAnswer = await buster.solveCaptchaImage(captchaImageLink)
10
11 // Fill the form with our solution
12 await tab.fill(".captchaForm",
13 { "captcha-answer": captchaAnswer },
14 { submit: true })
15 }
```

## Machine Readable Annotations

- Significantly reduce the effort to properly extract several entities
  - *Standard based*
  - *Limited amount of entities*
  - *Limited amount of relevant information*
- Content annotated by providers in any existing format
  - *Content-level*
    - *Microformats*
    - *Microdata*
    - *RDFa*
  - *Page-level*
    - *OpenGraph*
    - *JSON-LD*

# Content-Level Annotations

- Microformats

```
1 <p class="h-card">
2
3 Joe Bloggs
4 joebloggs@example.c
5 17 Austerstræti
6 Reykjavik
7 Iceland
8 </p>
```

- Microdata

```
1 <p itemscope itemprop="Person" itemtype="http://schema.org/Person">
2 Christopher Froome was sponsored by
3
4 Sky in the
5 </p>
```

- RDFa

```
1 <p vocab="http://schema.org/" typeof="Person">
2 Christopher Froome was sponsored by
3
4 Sky in the
5 </p>
```

# Page-Level Annotations

- OpenGraph

```
1 <head prefix="og: http://ogp.me/ns# fb: http://ogp.me/ns/fb# ">
2 <meta property="fb:app_id" content="45646845135" />
3 <meta property="og:url" content="http://example.com/profile/tom_h
4 <meta property="og:type" content="profile" />
5 <meta property="og:title" content="Tom Hanks" />
6 <meta property="og:description" content="Tom Hanks profile" />
7 <meta property="og:image" content="http://example.com/profile/hanks
```

- JSON-LD

```
1 <script type="application/ld+json">
2 {
3 "@context": "http://schema.org/",
4 "@type": "Person",
5 "name": "Christopher Froome",
6 "sponsor":
7 {
8 "@type": "Organization",
9 "name": "Sky",
10 "url": "http://www.skysports.com/"
11 }
12 }
13 </script>
```

## Machine Readable Annotations - Example

- BeautifulSoup

```
1 url="https://en.wikipedia.org/wiki/Web_scraping"
2 html_content = requests.get(url).text
3 soup = BeautifulSoup(html_content, "html.parser")
4 p = soup.find('script', {'type': 'application/ld+json'})
5 print(p.contents)

1 [{"@context": "https://schema.org", "@type": "Article", "name": "Web scrapin
```

- extract

```
1 import extract
2 import requests
3 import pprint
4 pp = pprint.PrettyPrinter(indent=2)
5 data = extract.extract(html_content)
6 pp.pprint(data)

1 { 'json-ld': [{ '@context': 'https://schema.org',
2 '@type': 'Article',
3 'author': { '@type': 'Organization',
4 'name': 'Contributors to Wikimedia projects',
5 'dateModified': '2019-12-16T19:58:14Z',
6 'datePublished': '2005-09-17T18:57:30Z',
7 'headline': 'data scraping used for extracting data from '
8 'websites',
```

## Other "Sources" for Scraping

- JavaScript variables

– `<script>` tag on the page

```
1 from selenium.webdriver import Firefox
2 from selenium.webdriver.firefox.options import Options
3 opts = Options()
4 browser = Firefox(options=opts)
5 browser.get('https://www.bbc.com/news/world-51235555')
6
7 browser.execute_script('return config;')
8
9 browser.close()

1 {
2 'headline': 'Space cookies: First food baked in space by astronauts',
3 'iStats_counter_name': 'news.world.story.51235555.page',
4 'language': 'en-gb',
5 'last_updated': {'date': '2020-01-24 12:26:02',
6 'timezone': 'Europe/London',
7 'timezone_type': 3},
8 'length': 2994,
9 'mediaType': 'video',
10 ...
11 'section': {'id': '99115',
12 'name': 'World',
13 'uri': '/news/world',
14 'urlIdentifier': '/news/world'},
```

## Other "Sources" for Scraping (cont.)

- API calls, XHRs
  - *website loads data dynamically*
  - *it typically uses XMLHttpRequests (XHRs)*

```
1 url = 'https://aukro.cz/backend/api/offers/search?page=0&size=60&sort='
2 data = {'text': 'notebook', "splitGroupKey": "listing", "splitGroupValue": "
3 headers = { 'Content-Type': 'application/json' }
4 requests.post(url, json = data, headers = headers).json()

1 ...
2 'content': [{ 'itemId': 121212121212,
3 'itemName': 'Notebook',
4 'catalog': { 'id': 1, 'name': 'Elektronika' },
5 ...
```

## Scraping Strategy

- Scrape Websites Without Being Blocked
  - *Slow down the scraping*
    - *too many request*
    - *too fast form submit etc.*
  - *Use proxy servers*
  - *Apply different scraping patterns*
  - *Look like a human*
    - *Use (switch) correct user-agents*
    - *Handle cookies*
  - *Be careful of honeypot traps*
  - *Use headless or real browser*