

Web 2.0

Lecture 7: HTTP/2

Ing. Milan Dojčinovski, Ph.D.

milan.dojcinovski@fit.cvut.cz • @m1ci • <http://dojchinovski.mk>



Czech Technical University in Prague
Faculty of Information Technologies • Software and Web Engineering • <https://courses.fit.cvut.cz/MI-W20/>



Evropský sociální fond
Praha & EU: Investujeme do vaší budoucnosti

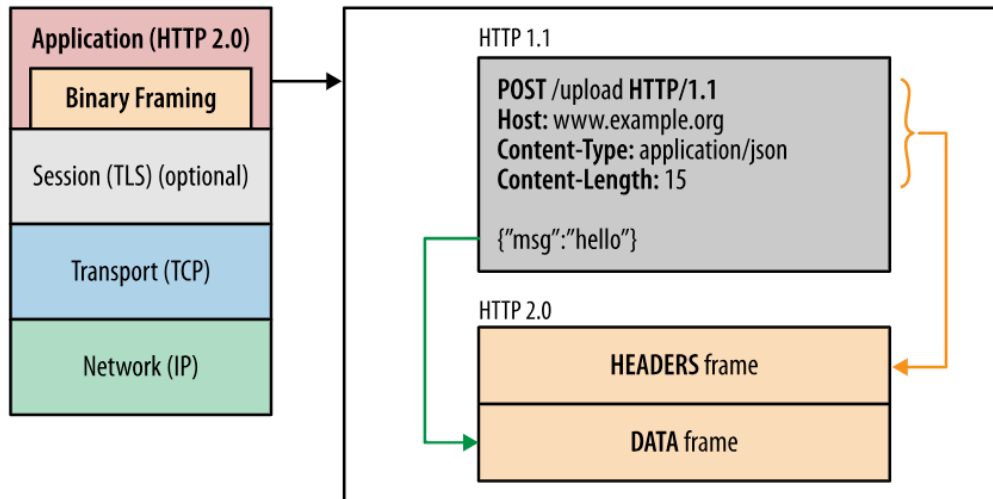
Modified: Mon Apr 12 2021, 12:42:32
Humla v0.3

Overview

- Developed from SPDY (2009) experimental protocol by Google
 - May 2015: RFC 7540 (HTTP/2) and RFC 7541 (HPACK)
 - HTTP/2 standards extend (not replace) the previous HTTP standards
- HTTP/1.x limitations
 - HTTP/1.x clients need to use multiple connections to achieve concurrency
 - unnecessary network traffic – request and response headers not compressed
 - no effective resource prioritization
- Primary goals
 - Reduction of latency
 - enabling full request and response multiplexing
 - minimize protocol overhead via efficient compression of HTTP header fields
 - support for request prioritization and server push
- HTTP/2 does not modify application semantics of HTTP
 - HTTP methods, URLs, header fields are the same
 - HTTP/2 modifies **how data is formatted and transported** in communication
- Literature and source
 - I. Grigorik: *High Performance Browser Networking*, O'Reilly Media, Inc. 2013.
ISBN: 9781449344757 [🔗](#)

Binary Framing Layer

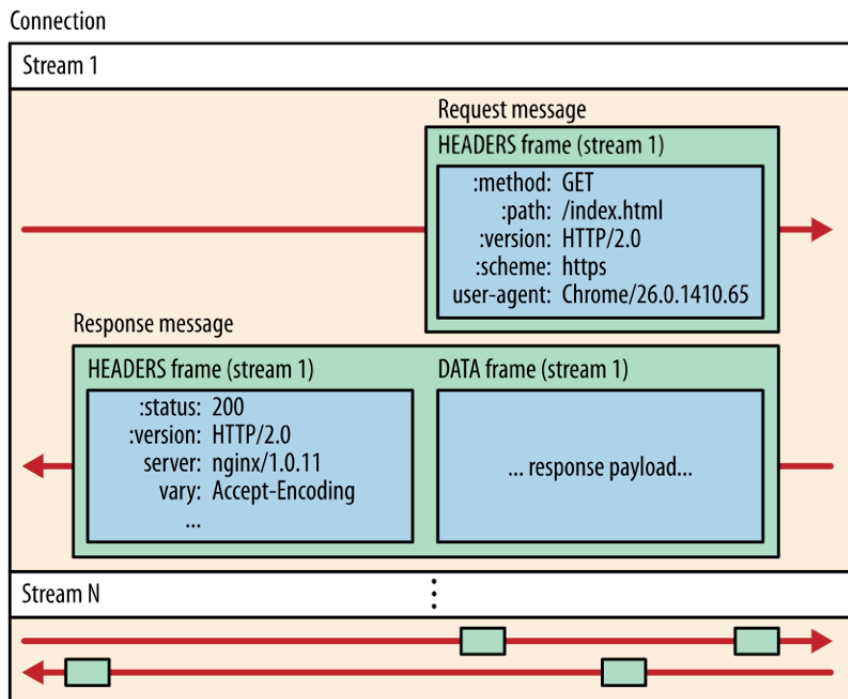
- Binary framing layer
 - defines how HTTP messages are encapsulated and transferred
 - communication is split into messages and frames in binary format



HTTP/2 Communication


- Data exchange between the client and server
 - break down of the communication into frames
 - frames are mapped to messages that belong to a particular stream
 - communication is **multiplexed within a single TCP connection**.
- Stream
 - bi-directional flow of bytes in a connection
 - may carry one or more messages
 - may have a priority
- Message
 - a sequence of frames
 - it maps to logical request or response message
- Frame
 - the smallest unit of communication
 - each has a frame header which identifies a stream to which it belongs.

Streams, messages, and frames




Request and response multiplexing

- Parallel requests
 - HTTP/1.x can use HTTP pipelining; they open multiple connections
 - browser typically opens up six connections
 - HTTP/2 allows full request and response multiplexing
 - Allows for parallel in-flight streams



- Performance benefits
 - Interleave requests/responses in parallel without blocking any one.
 - Deliver lower page load times by eliminating unnecessary latency
 - Improve utilization of available network capacity

HTTP/1.x Optimization



One connection

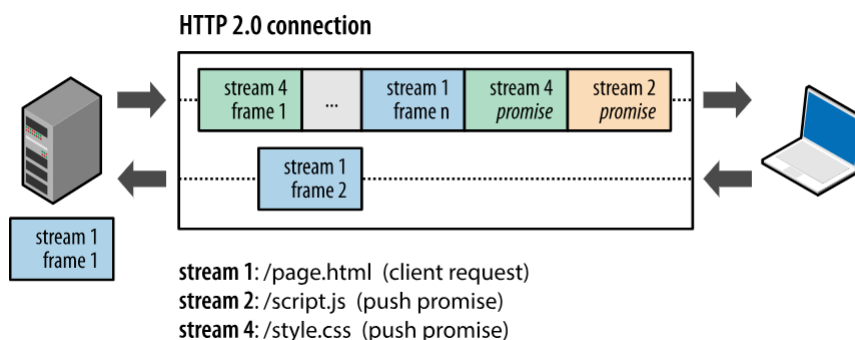
- Multiplexing allows for:
 - *all connections are persisted*
 - *only one connection required per origin*
- Advantages
 - *significant reduction of the overall protocol overhead*
 - *use of fewer connections reduces the memory and processing footprint along the full connection path (client, intermediaries, origin servers)*
 - *reduces operational costs and improves network capacity*
 - *Improves performance of HTTPS deployments*
 - *fewer expensive TLS handshakes*
 - *better session reuse*
 - *overall reduction in required client and server resources*

Flow control

- Prevent sender from receiving data it does not want
 - Receiver is busy or under heavy load
 - Receiver can allocate fixed amount of resources for particular stream
- Examples
 - Client request a video stream; a user pauses the stream
 - the client wants to pause the stream delivery to avoid buffering
 - A proxy server has a fast downstream and slow upstream
 - the proxy server can control how quickly the downstream delivers data to match the speed of upstream
 - better control of resource usage
 - Similar problems as in TCP control
- Flow control
 - Sender and receiver both advertise stream flow control window in bytes
 - = the size of the available buffer space to hold the incoming data
 - exchanged by special **SETTINGS** and **WINDOW_UPDATE** frames
 - Flow control is hop-by-hop, not end-to-end
 - an intermediary can set its own flow control

Server push

- Ability to send multiple responses for a single request
 - A response to the request is sent back
 - Additional resources can be pushed without client requesting them
 - Hypertext – "server knows what the client will need"




- Similar to resource inlining
 - A resource is pushed to the client in HTML/CSS resource
- Performance benefits
 - Cached by the client, reused across pages, multiplexed, declined by the client

Push promise

- **PUSH_PROMISE** frames
 - A signal that the server intends to push resources to the client
 - The client needs to know which resources the server intends to push to avoid creating duplicate requests for these resources.
- After the client receives **PUSH_PROMISE**
 - it may decline the stream (via **RST_STREAM** frame)
 - For example, when the resource is already in the cache
 - As for inline resources, this is not possible, the client always receives them
 - it can limit the number of concurrently pushed streams
 - it can adjust the initial flow control window to control how much data is pushed when the stream is first opened
 - it can disable server push entirely
- pushed resources must obey the same-origin policy

Stream Prioritization

- Purpose
 - Messages split into **frames** which are delivered in multiplexed **streams**
 - The order in which frames are delivered is important for a good performance
 - Client can define stream prioritization
 - optimizations in the browser; change prioritization based on user interaction
- Streams' weights and dependency
 - Each stream can be assigned an integer weight between 1 and 256.
 - Each stream may be given an explicit dependency on another stream.



Stream Prioritization (cont.)

- Dependency
 - referencing the unique identifier of another stream as its parent
 - if the identifier is omitted the stream is dependent on the "root stream"
 - The parent stream should be allocated resources ahead of its dependencies.
 - "Please process and deliver response D before response C"
- Weights
 - Sibling streams have resources allocated as per their weights
 - Example
 - Sum all the weights: $4 + 12 = 16$
 - Divide each stream weight by the total weight: $A = 12/16$, $B = 4/16$
 - Stream A receives $3/4$ and stream B receives $1/4$ of available resources;

Header compression

- Purpose
 - Each HTTP request/response contains a set of headers (metadata)
 - HTTP/1.x – metadata sent as plain text, adds 500-800 bytes per transfer
- HTTP/2 provides
 - Request and response metadata are compressed using HPACK format
 - header fields encoded via a static Huffman code – reduces size
 - client and server maintain an **indexed list of previously seen header fields**

